

Artificial Intelligence Foundation - **JC3001**

Lecture 11: Search IV: Adversarial Search II

Prof. Aladdin Ayesh (aladdin.ayesh@abdn.ac.uk)

Dr. Binod Bhattacharai (binod.bhattacharai@abdn.ac.uk)

Dr. Gideon Ogunniye, (g.ogunniye@abdn.ac.uk)

September 2025



UNIVERSITY OF
ABERDEEN



Material adapted from:
Russell and Norvig (AIMA Book): Chapter 5
Russell and Norvig (AIMA Book): Chapter 17/18 (17.1/18.2)
Shoham and Leyton-Brown (Game Theory)

- Part 1: Introduction
 - ① Introduction to AI ✓
 - ② Agents ✓
- Part 2: Problem-solving
 - ① Search 1: Uninformed Search ✓
 - ② Search 2: Heuristic Search ✓
 - ③ Search 3: Local Search ✓
 - ④ **Search 4: Adversarial Search**
- Part 3: Reasoning and Uncertainty
 - ① Reasoning 1: Constraint Satisfaction
 - ② Reasoning 2: Logic and Inference
 - ③ Probabilistic Reasoning 1: BNs
 - ④ Probabilistic Reasoning 2: HMMs
- Part 4: Planning
 - ① Planning 1: Intro and Formalism
 - ② Planning 2: Algos and Heuristics
 - ③ Planning 3: Hierarchical Planning
 - ④ Planning 4: Stochastic Planning
- Part 5: Learning
 - ① Learning 1: Intro to ML
 - ② Learning 2: Regression
 - ③ Learning 3: Neural Networks
 - ④ Learning 4: Reinforcement Learning
- Part 6: Conclusion
 - ① Ethical Issues in AI
 - ② Conclusions and Discussion

- Games ✓
- Non-deterministic search
- Adversarial Search
 - MinMax Search
 - Alpha-Beta Pruning
 - Expectimax



Outline

1 Non-deterministic Search

► Non-deterministic Search

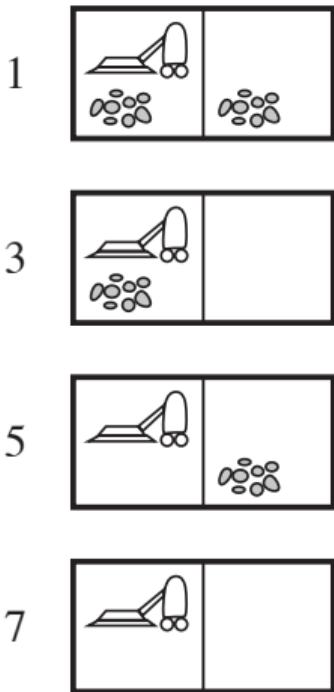
► Adversarial Search

- We have seen search algorithms for environments that are **deterministic** and **fully-observable**
- When either, or both of these assumptions cannot be made, an agent needs to account for **percepts** in every move to narrow down the states where it is
- A solution can no longer be a flat **sequence** of actions, but rather a **contingency plan** (a.k.a. **strategy**)

- Recall the vacuum world, but with a non-deterministic suck action:
 - When applied to a dirty square, the action cleans the current square, but sometimes also **cleans an adjacent square**
 - When applied to a clean square the action sometimes dirties the square
- This requires
 - Different **transition model**
 - Different **solution concept**

Non-deterministic Transitions

1 Non-deterministic Search



- Result from a state now becomes Results:

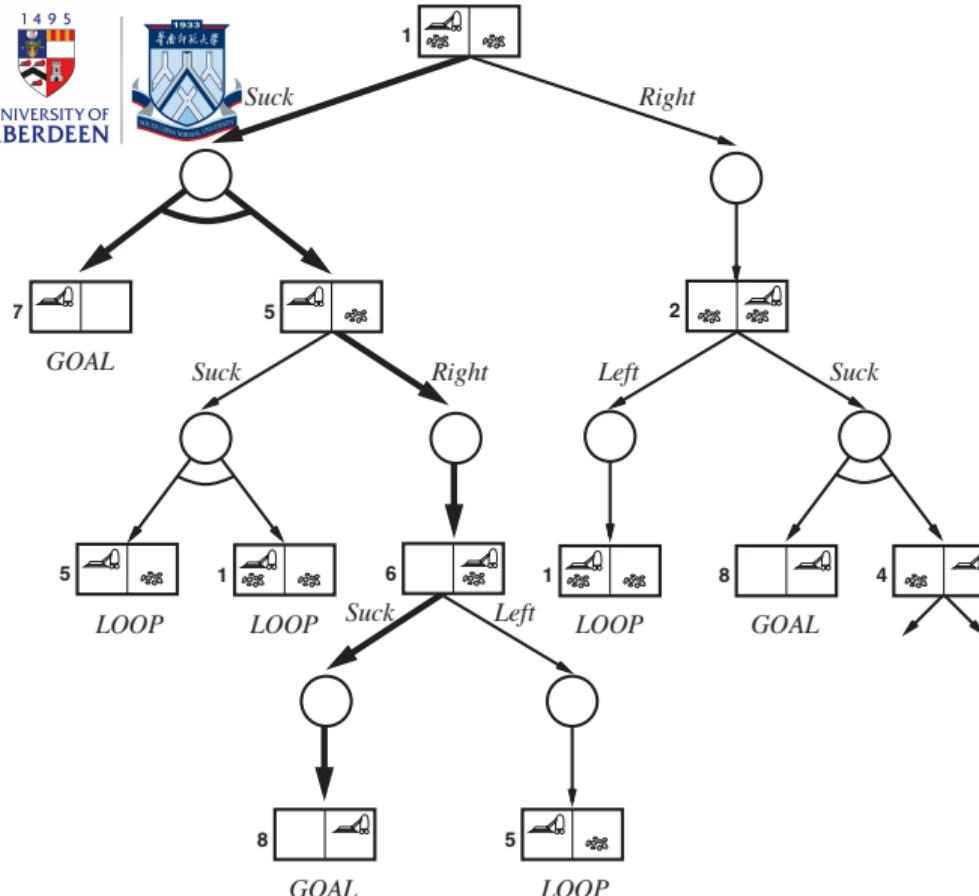
$$\text{Results}(1, \text{suck}) = \{5, 7\}$$

Suck applied to state 1 results in the agent arriving either in states 5 or 7

- Plan is now a contingency plan (a tree):

[Suck, if State = 5 then [Right, Suck] else []]

- In deterministic environments, branching only occurs due to agent's choice (OR Nodes)
- In non-deterministic environments, the environment's choice must also be taken into account (AND Nodes)
- Solution is a subtree of the AND-OR tree that:
 - Has a goal node at every leaf
 - Specifies an action at each OR node
 - Includes every outcome branch of its AND nodes



AND-OR Trees

1 Non-deterministic Search

First two levels of AND-OR
Search Tree for erratic
vacuum

- Or-Search: similar to regular search (agent has choice)
- And-Search: agent searching against the environment
- Notice failure condition on Or-Search

```

function And-Or-Graph-Search(problem)
    return Or-Search(problem.Initial-State, problem, [])

function Or-Search(state, problem, path)
    if problem.Goal-Test(state) then return []
    if state is on path then return failure
    for each action in problem.Actions(state) do
        plan  $\leftarrow$  And-Search(Results(state, action), problem, [state|path])
        if plan  $\neq$  failure then return [action|plan]
    return failure

function And-Search(states, problem, path)
    for each si in states do
        plan  $\leftarrow$  Or-Search(si, problem, path)
        if plan = failure then return failure
    return [if s1 then plan1 else if s2 then plan2 else . . . if sn-1 then plann-1 else plann]

```



Outline

2 Adversarial Search

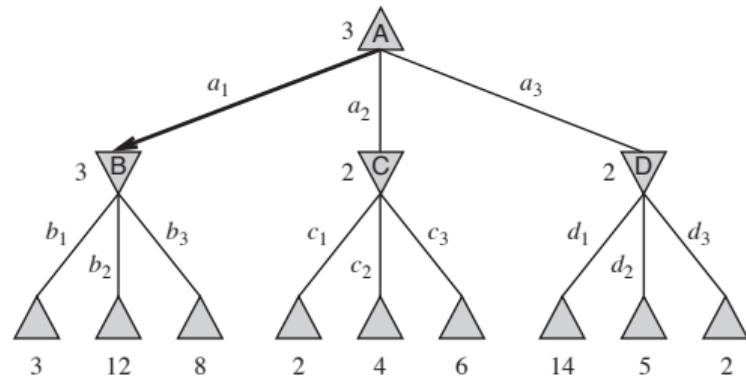
- ▶ Non-deterministic Search
- ▶ Adversarial Search

Adversarial Optimal Decisions

2 Adversarial Search

MAX

MIN



- Each move we (MAX) makes has a response from MIN
- So the plan we look for is a **contingent strategy** with moves for:
 - The initial state; and
 - Every possible response from MIN
- We can calculate the value of the game recursively, assuming both players play optimally using the **Minimax** value

$$\text{Minimax}(s) = \begin{cases} \text{Utility}(s) & \text{if Terminal-Test}(s) \\ \max_{a \in \text{Actions}(s)} \text{Minimax}(\text{Result}(s, a)) & \text{if Player}(s) = \text{Max} \\ \min_{a \in \text{Actions}(s)} \text{Minimax}(\text{Result}(s, a)) & \text{if Player}(s) = \text{Min} \end{cases}$$



UNIVERSITY OF
ABERDEEN



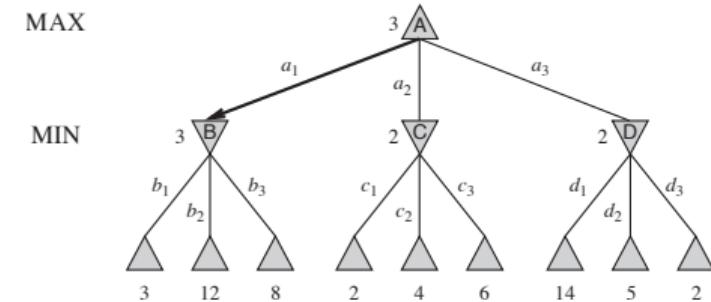
Minimax Algorithm

2 Adversarial Search

```
1: function Minimax-Decision(state) returns an action
2:   return  $\arg \max_{a \in \text{Actions}(s)} \text{Min-Value}(\text{Result}(s, a))$ 
```

```
3: function Max-Value(state) returns a utility value
4:   if Terminal-Test(state) then return Utility(state)
5:    $v \leftarrow -\infty$ 
6:   for each a in Actions(state) do
7:      $v \leftarrow \text{Max}(v, \text{Min-Value}(\text{Result}(s, a)))$ 
8:   return v
```

```
9: function Min-Value(state) returns a utility value
10:  if Terminal-Test(state) then return Utility(state)
11:   $v \leftarrow \infty$ 
12:  for each a in Actions(state) do
13:     $v \leftarrow \text{Min}(v, \text{Max-Value}(\text{Result}(s, a)))$ 
14:  return v
```



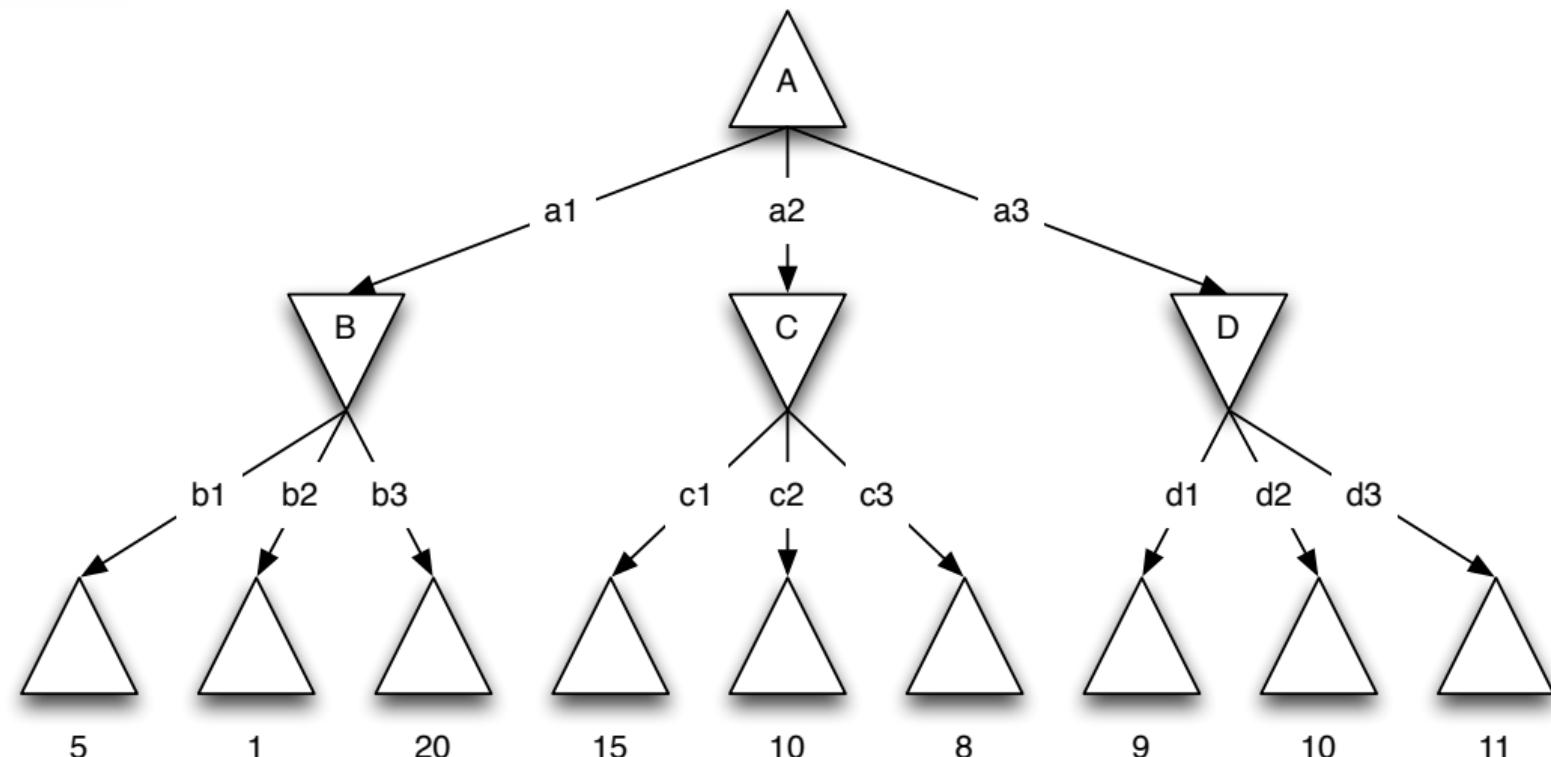


UNIVERSITY OF
ABERDEEN



Minimax Value Question

2 Adversarial Search



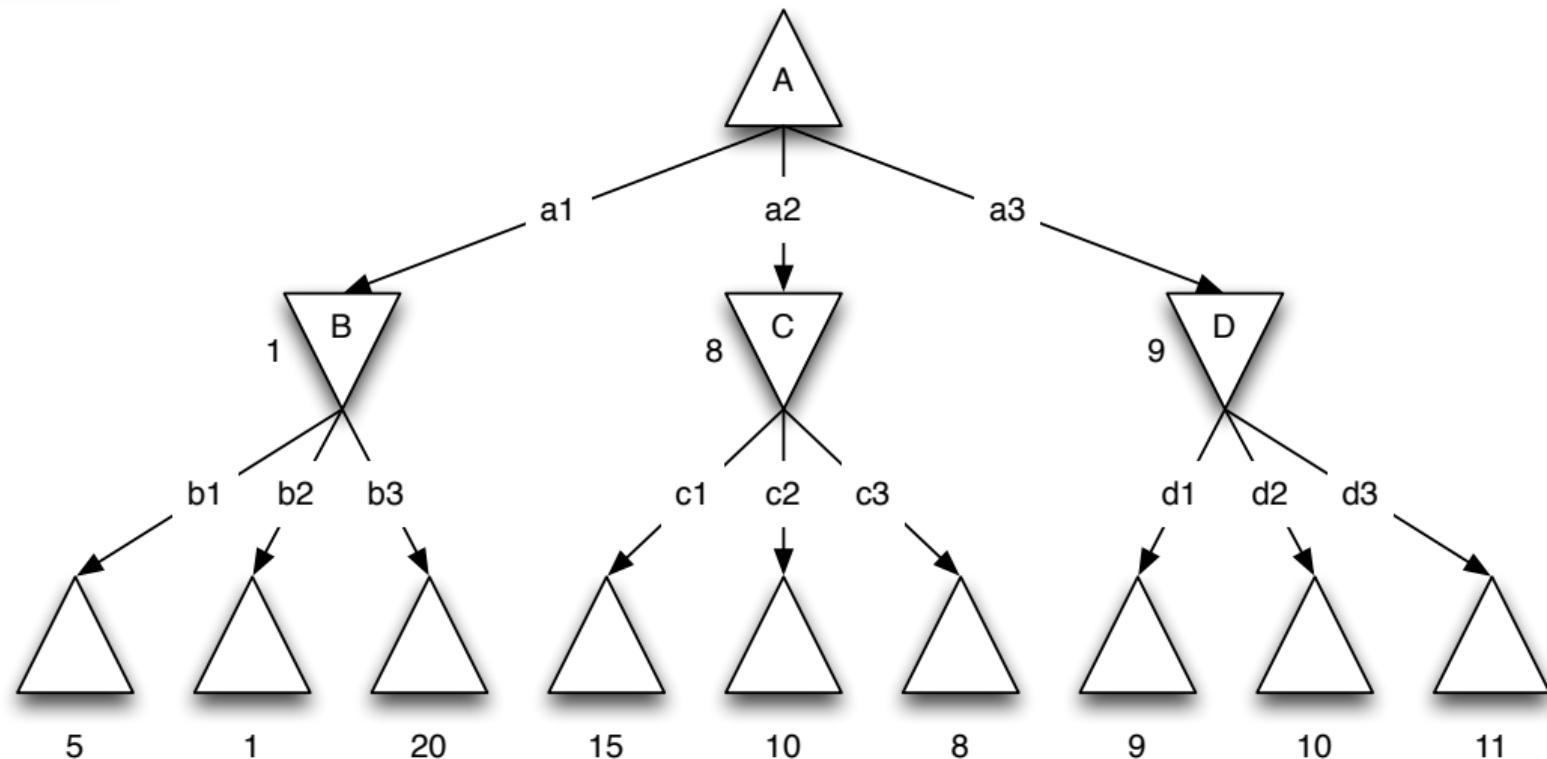


UNIVERSITY OF
ABERDEEN



Minimax Value Question

2 Adversarial Search



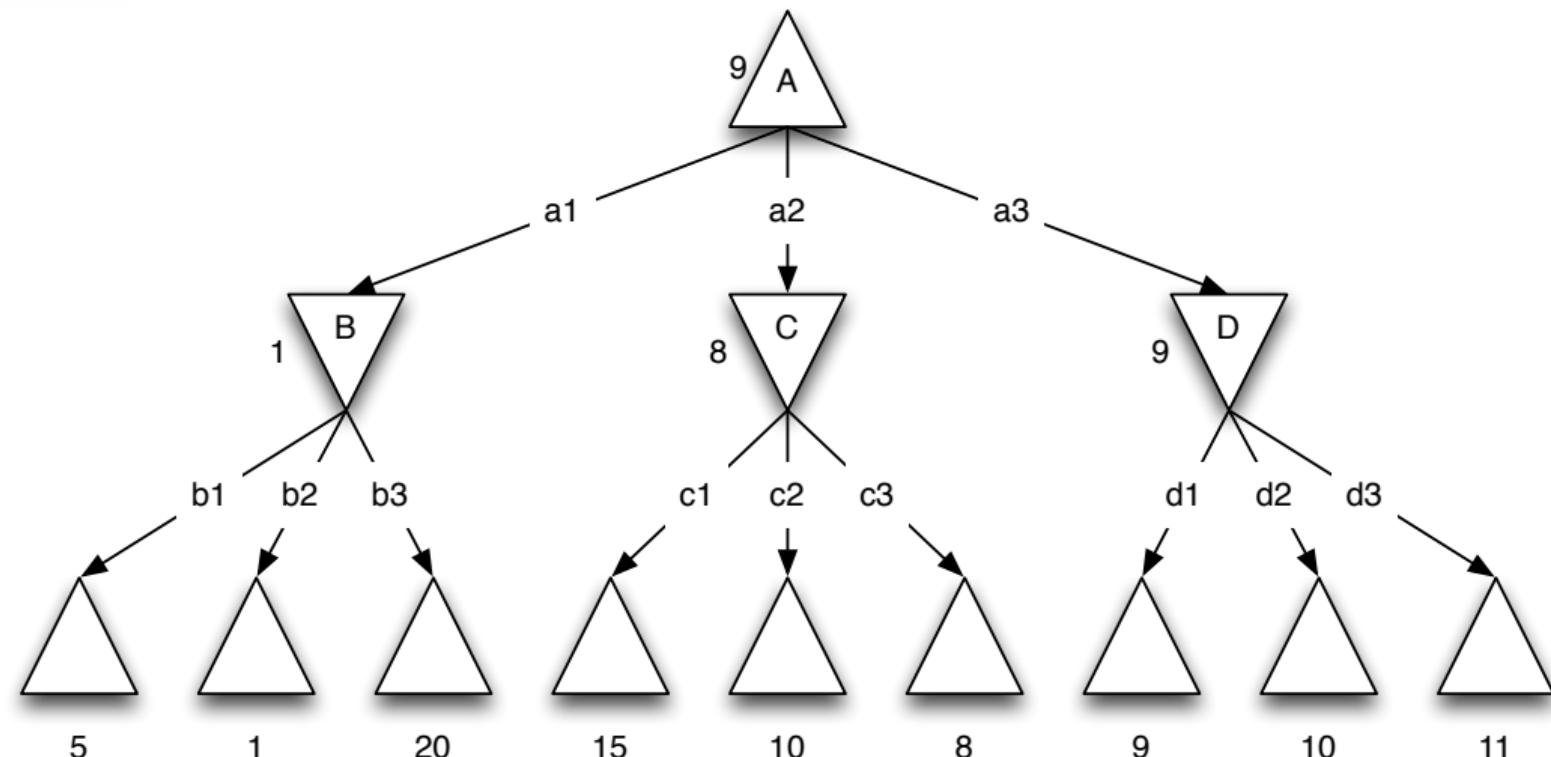


UNIVERSITY OF
ABERDEEN



Minimax Value Question

2 Adversarial Search



```
1: function Minimax-Decision(state) returns an action
2:   return  $\arg \max_{a \in \text{Actions}(s)} \text{Min-Value}(\text{Result}(s, a))$ 
```

```
3: function Max-Value(state) returns a utility value
4:   if Terminal-Test(state) then return Utility(state)
5:    $v \leftarrow -\infty$ 
6:   for each a in Actions(state) do
7:      $v \leftarrow \text{Max}(v, \text{Min-Value}(\text{Result}(s, a)))$ 
8:   return v
```

```
9: function Min-Value(state) returns a utility value
10:  if Terminal-Test(state) then return Utility(state)
11:   $v \leftarrow \infty$ 
12:  for each a in Actions(state) do
13:     $v \leftarrow \text{Min}(v, \text{Max-Value}(\text{Result}(s, a)))$ 
14:  return v
```

- Time Complexity
- Space Complexity

```
1: function Minimax-Decision(state) returns an action
2:   return  $\arg \max_{a \in \text{Actions}(s)} \text{Min-Value}(\text{Result}(s, a))$ 
```

```
3: function Max-Value(state) returns a utility value
4:   if Terminal-Test(state) then return Utility(state)
5:    $v \leftarrow -\infty$ 
6:   for each a in Actions(state) do
7:      $v \leftarrow \text{Max}(v, \text{Min-Value}(\text{Result}(\text{state}, a)))$ 
8:   return v
```

```
9: function Min-Value(state) returns a utility value
10:  if Terminal-Test(state) then return Utility(state)
11:   $v \leftarrow \infty$ 
12:  for each a in Actions(state) do
13:     $v \leftarrow \text{Min}(v, \text{Max-Value}(\text{Result}(\text{state}, a)))$ 
14:  return v
```

- Time Complexity
 $O(b^m)$
- Space Complexity

```
1: function Minimax-Decision(state) returns an action
2:   return  $\arg \max_{a \in \text{Actions}(s)} \text{Min-Value}(\text{Result}(s, a))$ 
```

```
3: function Max-Value(state) returns a utility value
4:   if Terminal-Test(state) then return Utility(state)
5:    $v \leftarrow -\infty$ 
6:   for each a in Actions(state) do
7:      $v \leftarrow \text{Max}(v, \text{Min-Value}(\text{Result}(\text{state}, a)))$ 
8:   return v
```

```
9: function Min-Value(state) returns a utility value
10:  if Terminal-Test(state) then return Utility(state)
11:   $v \leftarrow \infty$ 
12:  for each a in Actions(state) do
13:     $v \leftarrow \text{Min}(v, \text{Max-Value}(\text{Result}(\text{state}, a)))$ 
14:  return v
```

- Time Complexity
 $O(b^m)$
- Space Complexity
 $O(bm)$

```
1: function Minimax-Decision(state) returns an action
2:   return  $\arg \max_{a \in \text{Actions}(s)} \text{Min-Value}(\text{Result}(s, a))$ 
```

```
3: function Max-Value(state) returns a utility value
4:   if Terminal-Test(state) then return Utility(state)
5:    $v \leftarrow -\infty$ 
6:   for each a in Actions(state) do
7:      $v \leftarrow \text{Max}(v, \text{Min-Value}(\text{Result}(\text{state}, a)))$ 
8:   return v
```

```
9: function Min-Value(state) returns a utility value
10:  if Terminal-Test(state) then return Utility(state)
11:   $v \leftarrow \infty$ 
12:  for each a in Actions(state) do
13:     $v \leftarrow \text{Min}(v, \text{Max-Value}(\text{Result}(\text{state}, a)))$ 
14:  return v
```

- Time Complexity
 $O(b^m)$
- Space Complexity
 $O(bm)$
- Chess, on average:
 $b = 30 m = 40$

- Reducing complexity of b^m
 - Reduce branching factor (b)?
 - Reduce maximum search depth (m)?
 - Searching in a graph rather than a tree?



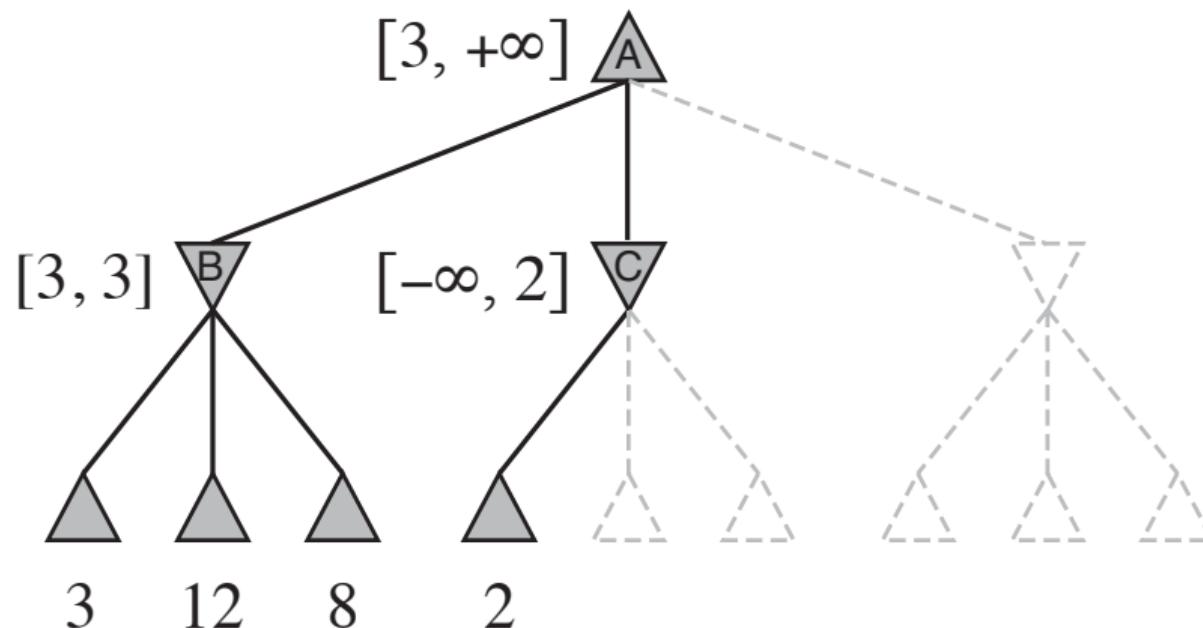
UNIVERSITY OF
ABERDEEN



1933
香港中文大學
The Chinese University of Hong Kong

Reducing Branching Factor

2 Adversarial Search



- Alpha-Beta Pruning
 - Evaluate which nodes/branches would not affect MIN/MAX's decision
 - Based on keeping track of two parameters:
 - α - value of the best (highest) choice we have in MAX's path
 - β - value of the best (lowest) choice we have in MIN's path
- Updates these values as one goes along the tree



UNIVERSITY OF
ABERDEEN



Alpha-Beta Pruning

2 Adversarial Search

```
1: function Minimax-Decision(state) returns an action
2:    $v \leftarrow \text{Min-Value}(\textit{state}, -\infty, +\infty)$ 
3:   return the action in Actions(state) with value  $v$ 
```

```
4: function Max-Value(state) returns a utility value
5:   if Terminal-Test(state) then return Utility(state)
6:    $v \leftarrow -\infty$ 
7:   for each a in Actions(state) do
8:      $v \leftarrow \text{Max}(v, \text{Min-Value}(\text{Result}(\textit{state}, a), \alpha, \beta))$ 
```

```
11: return  $v$ 
```

```
12: function Min-Value(state) returns a utility value
13:   if Terminal-Test(state) then return Utility(state)
14:    $v \leftarrow \infty$ 
15:   for each a in Actions(state) do
16:      $v \leftarrow \text{Min}(v, \text{Max-Value}(\text{Result}(\textit{state}, a), \alpha, \beta))$ 
```

```
19: return  $v$ 
```

```

1: function Minimax-Decision(state) returns an action
2:    $v \leftarrow \text{Min-Value}(\textit{state}, -\infty, +\infty)$ 
3:   return the action in Actions(state) with value  $v$ 

```

```

4: function Max-Value(state) returns a utility value
5:   if Terminal-Test(state) then return Utility(state)
6:    $v \leftarrow -\infty$ 
7:   for each a in Actions(state) do
8:      $v \leftarrow \text{Max}(v, \text{Min-Value}(\text{Result}(\textit{state}, a), \alpha, \beta))$ 
10:     $\alpha \leftarrow \text{Max}(\alpha, v)$ 
11:   return  $v$ 

```

```

12: function Min-Value(state) returns a utility value
13:   if Terminal-Test(state) then return Utility(state)
14:    $v \leftarrow \infty$ 
15:   for each a in Actions(state) do
16:      $v \leftarrow \text{Min}(v, \text{Max-Value}(\text{Result}(\textit{state}, a), \alpha, \beta))$ 
18:      $\beta \leftarrow \text{Min}(\beta, v)$ 
19:   return  $v$ 

```



UNIVERSITY OF
ABERDEEN

Alpha-Beta Pruning

2 Adversarial Search

```
1: function Minimax-Decision(state) returns an action
2:    $v \leftarrow \text{Min-Value}(\textit{state}, -\infty, +\infty)$ 
3:   return the action in Actions(state) with value  $v$ 
```

```
4: function Max-Value(state) returns a utility value
5:   if Terminal-Test(state) then return Utility(state)
6:    $v \leftarrow -\infty$ 
7:   for each a in Actions(state) do
8:      $v \leftarrow \text{Max}(v, \text{Min-Value}(\text{Result}(\textit{state}, a), \alpha, \beta))$ 
9:     if  $v \geq \beta$  then return  $v$ 
10:     $\alpha \leftarrow \text{Max}(\alpha, v)$ 
11:   return  $v$ 
```

```
12: function Min-Value(state) returns a utility value
13:   if Terminal-Test(state) then return Utility(state)
14:    $v \leftarrow \infty$ 
15:   for each a in Actions(state) do
16:      $v \leftarrow \text{Min}(v, \text{Max-Value}(\text{Result}(\textit{state}, a), \alpha, \beta))$ 
17:     if  $v \leq \alpha$  then return  $v$ 
18:      $\beta \leftarrow \text{Min}(\beta, v)$ 
19:   return  $v$ 
```



UNIVERSITY OF
ABERDEEN

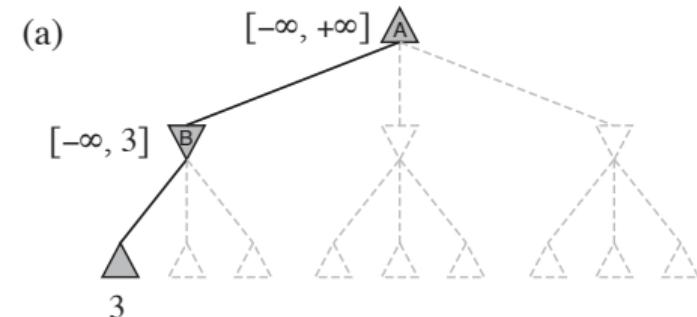
Alpha-Beta Pruning

2 Adversarial Search

```
1: function Minimax-Decision(state) returns an action
2:    $v \leftarrow \text{Min-Value}(state, -\infty, +\infty)$ 
3:   return the action in Actions(state) with value  $v$ 
```

```
4: function Max-Value(state) returns a utility value
5:   if Terminal-Test(state) then return Utility(state)
6:    $v \leftarrow -\infty$ 
7:   for each  $a$  in Actions(state) do
8:      $v \leftarrow \text{Max}(v, \text{Min-Value}(\text{Result}(state, a), \alpha, \beta))$ 
9:     if  $v \geq \beta$  then return  $v$ 
10:     $\alpha \leftarrow \text{Max}(\alpha, v)$ 
11:   return  $v$ 
```

```
12: function Min-Value(state) returns a utility value
13:   if Terminal-Test(state) then return Utility(state)
14:    $v \leftarrow \infty$ 
15:   for each  $a$  in Actions(state) do
16:      $v \leftarrow \text{Min}(v, \text{Max-Value}(\text{Result}(state, a), \alpha, \beta))$ 
17:     if  $v \leq \alpha$  then return  $v$ 
18:      $\beta \leftarrow \text{Min}(\beta, v)$ 
19:   return  $v$ 
```





UNIVERSITY OF
ABERDEEN

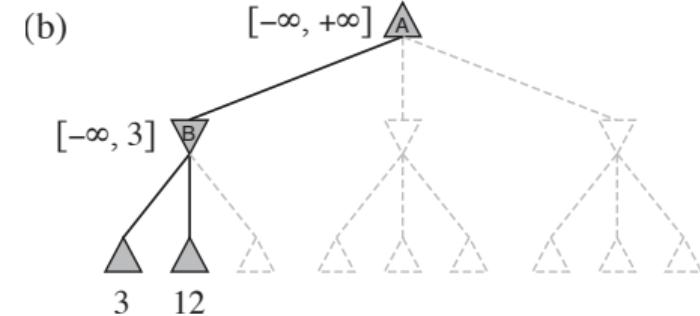
Alpha-Beta Pruning

2 Adversarial Search

```
1: function Minimax-Decision(state) returns an action
2:    $v \leftarrow \text{Min-Value}(state, -\infty, +\infty)$ 
3:   return the action in Actions(state) with value  $v$ 
```

```
4: function Max-Value(state) returns a utility value
5:   if Terminal-Test(state) then return Utility(state)
6:    $v \leftarrow -\infty$ 
7:   for each  $a$  in Actions(state) do
8:      $v \leftarrow \text{Max}(v, \text{Min-Value}(\text{Result}(state, a), \alpha, \beta))$ 
9:     if  $v \geq \beta$  then return  $v$ 
10:     $\alpha \leftarrow \text{Max}(\alpha, v)$ 
11:   return  $v$ 
```

```
12: function Min-Value(state) returns a utility value
13:   if Terminal-Test(state) then return Utility(state)
14:    $v \leftarrow \infty$ 
15:   for each  $a$  in Actions(state) do
16:      $v \leftarrow \text{Min}(v, \text{Max-Value}(\text{Result}(state, a), \alpha, \beta))$ 
17:     if  $v \leq \alpha$  then return  $v$ 
18:      $\beta \leftarrow \text{Min}(\beta, v)$ 
19:   return  $v$ 
```





UNIVERSITY OF
ABERDEEN

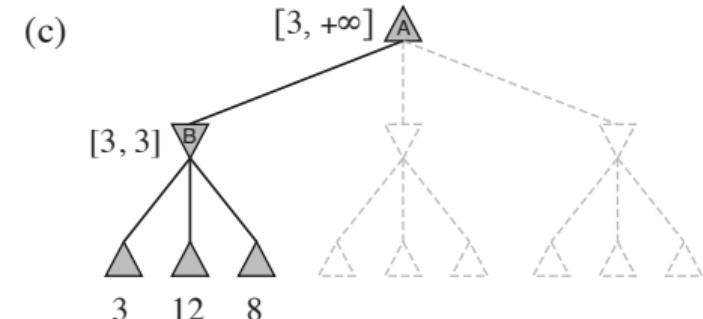
Alpha-Beta Pruning

2 Adversarial Search

```
1: function Minimax-Decision(state) returns an action
2:    $v \leftarrow \text{Min-Value}(state, -\infty, +\infty)$ 
3:   return the action in Actions(state) with value  $v$ 
```

```
4: function Max-Value(state) returns a utility value
5:   if Terminal-Test(state) then return Utility(state)
6:    $v \leftarrow -\infty$ 
7:   for each  $a$  in Actions(state) do
8:      $v \leftarrow \text{Max}(v, \text{Min-Value}(\text{Result}(state, a), \alpha, \beta))$ 
9:     if  $v \geq \beta$  then return  $v$ 
10:     $\alpha \leftarrow \text{Max}(\alpha, v)$ 
11:   return  $v$ 
```

```
12: function Min-Value(state) returns a utility value
13:   if Terminal-Test(state) then return Utility(state)
14:    $v \leftarrow \infty$ 
15:   for each  $a$  in Actions(state) do
16:      $v \leftarrow \text{Min}(v, \text{Max-Value}(\text{Result}(state, a), \alpha, \beta))$ 
17:     if  $v \leq \alpha$  then return  $v$ 
18:      $\beta \leftarrow \text{Min}(\beta, v)$ 
19:   return  $v$ 
```



```

1: function Minimax-Decision(state) returns an action
2:    $v \leftarrow \text{Min-Value}(\text{state}, -\infty, +\infty)$ 
3:   return the action in Actions(state) with value  $v$ 

```

```

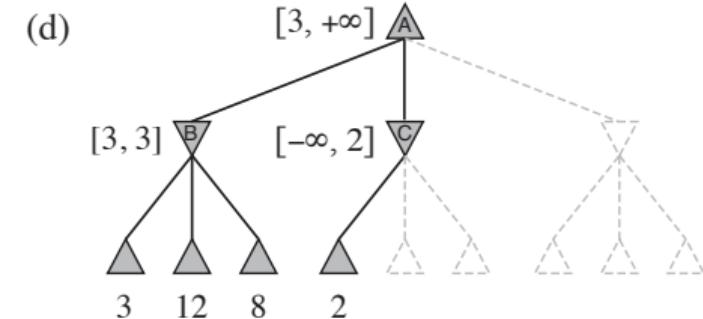
4: function Max-Value(state) returns a utility value
5:   if Terminal-Test(state) then return Utility(state)
6:    $v \leftarrow -\infty$ 
7:   for each a in Actions(state) do
8:      $v \leftarrow \text{Max}(v, \text{Min-Value}(\text{Result}(\text{state}, a), \alpha, \beta))$ 
9:     if  $v \geq \beta$  then return  $v$ 
10:     $\alpha \leftarrow \text{Max}(\alpha, v)$ 
11:   return  $v$ 

```

```

12: function Min-Value(state) returns a utility value
13:   if Terminal-Test(state) then return Utility(state)
14:    $v \leftarrow \infty$ 
15:   for each a in Actions(state) do
16:      $v \leftarrow \text{Min}(v, \text{Max-Value}(\text{Result}(\text{state}, a), \alpha, \beta))$ 
17:     if  $v \leq \alpha$  then return  $v$ 
18:      $\beta \leftarrow \text{Min}(\beta, v)$ 
19:   return  $v$ 

```





UNIVERSITY OF
ABERDEEN

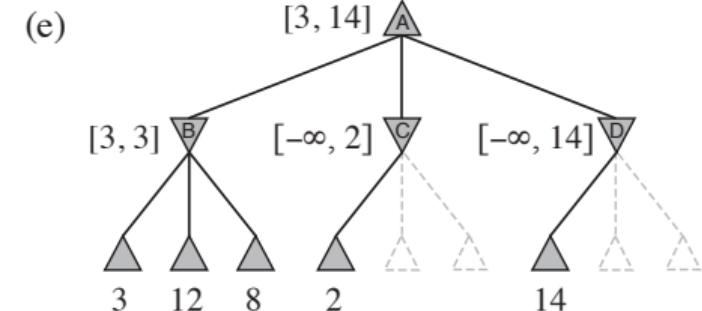
```
1: function Minimax-Decision(state) returns an action
2:    $v \leftarrow \text{Min-Value}(state, -\infty, +\infty)$ 
3:   return the action in Actions(state) with value  $v$ 
```

```
4: function Max-Value(state) returns a utility value
5:   if Terminal-Test(state) then return Utility(state)
6:    $v \leftarrow -\infty$ 
7:   for each  $a$  in Actions(state) do
8:      $v \leftarrow \text{Max}(v, \text{Min-Value}(\text{Result}(state, a), \alpha, \beta))$ 
9:     if  $v \geq \beta$  then return  $v$ 
10:     $\alpha \leftarrow \text{Max}(\alpha, v)$ 
11:  return  $v$ 
```

```
12: function Min-Value(state) returns a utility value
13:   if Terminal-Test(state) then return Utility(state)
14:    $v \leftarrow \infty$ 
15:   for each  $a$  in Actions(state) do
16:      $v \leftarrow \text{Min}(v, \text{Max-Value}(\text{Result}(state, a), \alpha, \beta))$ 
17:     if  $v \leq \alpha$  then return  $v$ 
18:      $\beta \leftarrow \text{Min}(\beta, v)$ 
19:  return  $v$ 
```

Alpha-Beta Pruning

2 Adversarial Search





UNIVERSITY OF
ABERDEEN

1: **function** Minimax-Decision(*state*) **returns** an action
2: $v \leftarrow \text{Min-Value}(\text{state}, -\infty, +\infty)$
3: **return** the *action* in Actions(*state*) with value v

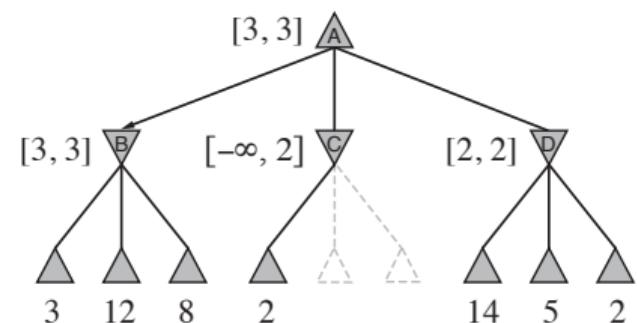
4: **function** Max-Value(*state*) **returns** a utility value
5: **if** Terminal-Test(*state*) **then return** Utility(*state*)
6: $v \leftarrow -\infty$
7: **for each** *a* in Actions(*state*) **do**
8: $v \leftarrow \text{Max}(v, \text{Min-Value}(\text{Result}(\text{state}, a), \alpha, \beta))$
9: **if** $v \geq \beta$ **then return** v
10: $\alpha \leftarrow \text{Max}(\alpha, v)$
11: **return** v

12: **function** Min-Value(*state*) **returns** a utility value
13: **if** Terminal-Test(*state*) **then return** Utility(*state*)
14: $v \leftarrow \infty$
15: **for each** *a* in Actions(*state*) **do**
16: $v \leftarrow \text{Min}(v, \text{Max-Value}(\text{Result}(\text{state}, a), \alpha, \beta))$
17: **if** $v \leq \alpha$ **then return** v
18: $\beta \leftarrow \text{Min}(\beta, v)$
19: **return** v

Alpha-Beta Pruning

2 Adversarial Search

(f)



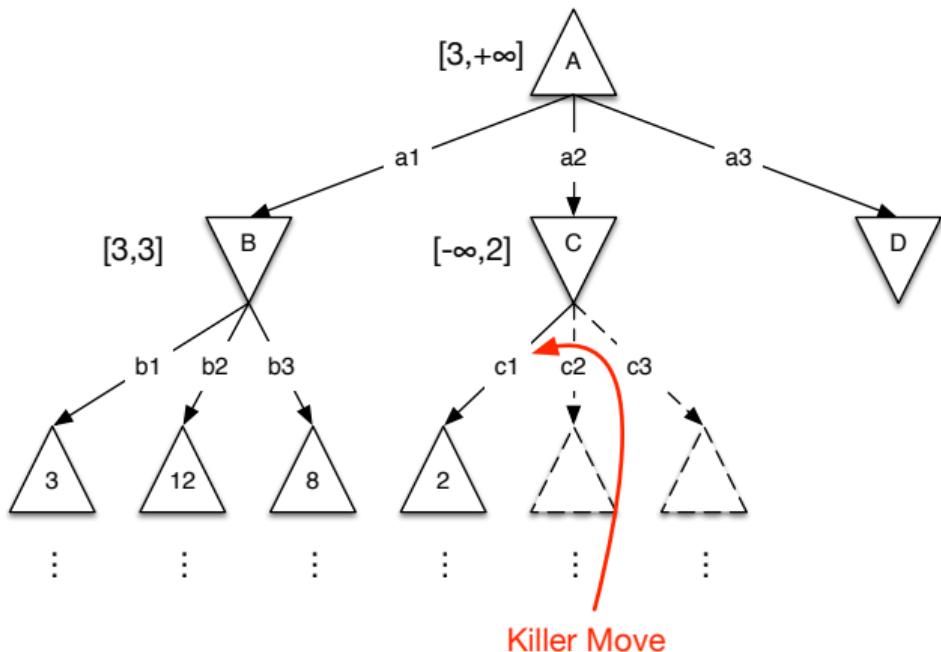
- Pruning is strongly affected by the ordering of the moves in the tree
 - A good ordering*, would enable us to prune many nodes
- Move ordering is often game-dependent knowledge (heuristic)
- Dynamic move-ordering (killer-move heuristic)

- Dynamic heuristic to determine a “good” ordering
- Search two plies ahead until Max (alt. Min) causes a beta (alt. alpha) cutoff
- The move that caused the cutoff is the killer move

Reducing D - Killer Move

2 Adversarial Search

- Dynamic heuristic to determine a “good” ordering
- Search two plies ahead until Max (alt. Min) causes a beta (alt. alpha) cutoff
- The move that caused the cutoff is the killer move





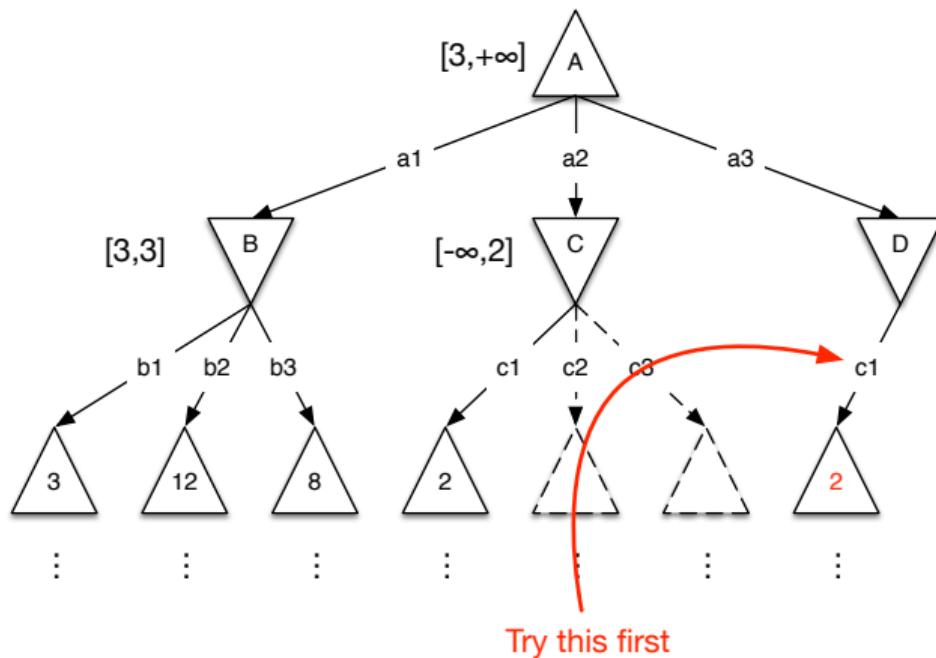
UNIVERSITY OF
ABERDEEN



Reducing D - Killer Move

2 Adversarial Search

- Dynamic heuristic to determine a “good” ordering
- Search two plies ahead until Max (alt. Min) causes a beta (alt. alpha) cutoff
- The move that caused the cutoff is the killer move



- Establish a cutoff depth and then estimate the value of future paths
 - Estimated value can be a heuristic (domain knowledge)
 - Values can be learned from previous games

$$\text{H-Minimax}(s, d) = \begin{cases} \text{Eval}(s) & \text{if } \text{Cutoff-Test}(s, d) \\ \max_{a \in \text{Actions}(s)} \text{H-Minimax}(\text{Result}(s, a), d + 1) & \text{if } \text{Player}(s) = \text{Max} \\ \min_{a \in \text{Actions}(s)} \text{H-Minimax}(\text{Result}(s, a), d + 1) & \text{if } \text{Player}(s) = \text{Min} \end{cases}$$

- Weighted linear function over features of a state

$$\text{Eval}(s) = w_1 f_1(s) + w_2 f_2(s) + \dots w_n f_n(s) = \sum_{i=1}^n w_i f_i(s)$$

- Example: Chess
Current state: pieces, and positions (structure)
- Example: Magic (Card Game)
Current state: Life Points, Cards in Play and Hand

- Evaluation function can also be learned (i.e. machine learning)
 - At cutoff we have a weighted linear function

$$\text{Eval}(s) = \sum_{i=1}^n w_i f_i(s)$$

- Record the features and actual value V at each evaluation

$$w_1, w_2, \dots, w_n, V$$

- Use supervised learning on weights w_i to approximate actual outcome

- As in non-adversarial search, many states will be revisited
- However, only recording visited states is not enough (since MIN can deviate in the future)
- Need to store actual loop paths (memory intensive)
 - Requires “caching” strategy

- Many tabletop games, and most simulated games include an element of chance
 - e.g., Backgammon
- Outcome of agent choices is not deterministic
 - Games must take into account multiple outcomes for the player
- Solution: weight outcomes by their probability
 - Expected value



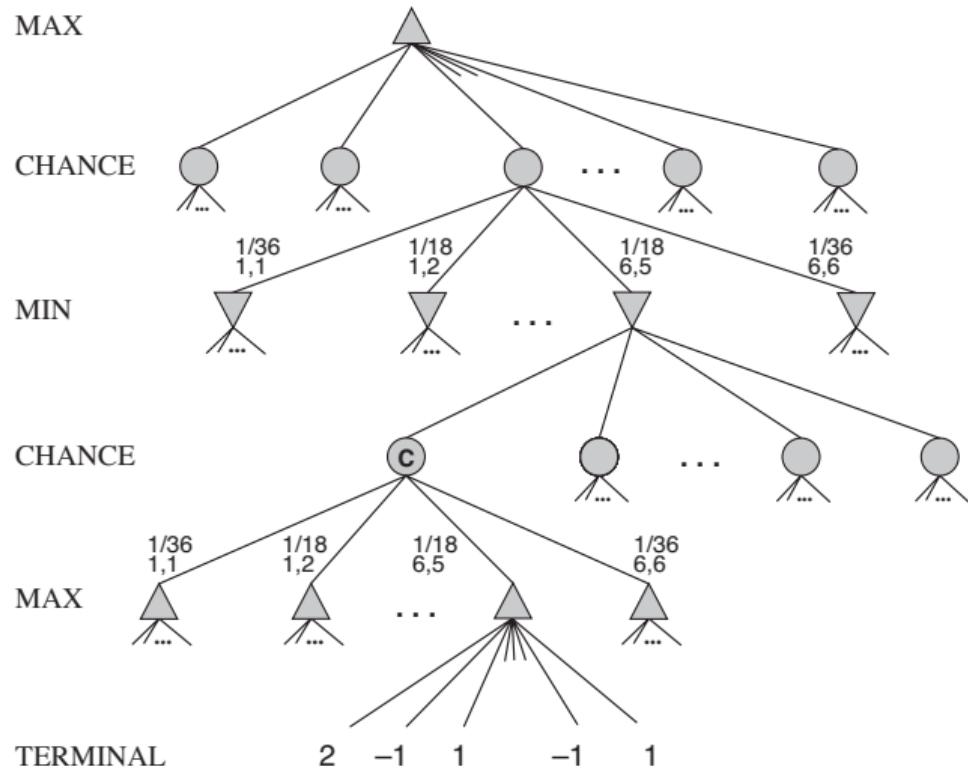
UNIVERSITY OF
ABERDEEN



Instead of the deterministic utilities, account for stochastic utilities

Expectiminimax

2 Adversarial Search



$$\text{ExpectMinimax}(s, d) = \begin{cases} \text{Utility}(s) & \text{if Terminal-Test}(s, d) \\ \max_{a \in \text{Actions}(s)} \text{ExpectMinimax}(\text{Result}(s, a)) & \text{if Player}(s) = \text{Max} \\ \min_{a \in \text{Actions}(s)} \text{ExpectMinimax}(\text{Result}(s, a)) & \text{if Player}(s) = \text{Min} \\ \sum_r \mathbb{P}[r] \text{ExpectMinimax}(\text{Result}(s, a)) & \text{if Player}(s) = \text{Chance} \end{cases}$$

- Since the outcomes of all actions are stochastic, instead of computing all expected utilities, we can sample game outcomes
 - Simulate games but choose **one random outcome** per player choice
 - Monte Carlo Rollout

- We've seen how to model games:
 - Building a game-tree
 - Finding the optimal moves in fully observable settings
- There are many other types of games
 - Multiplayer games
 - Stochastic Games
 - Partially-observable games



UNIVERSITY OF
ABERDEEN



Any Questions.