

Artificial Intelligence Foundation - JC3001

Lecture 15: Constraint Satisfaction Problems III

Prof. Aladdin Ayesh (aladdin.ayesh@abdn.ac.uk)

Dr. Binod Bhattacharai (binod.bhattacharai@abdn.ac.uk)

Dr. Gideon Ogunniye, (g.ogunniye@abdn.ac.uk)

September 2025



UNIVERSITY OF
ABERDEEN



Material adapted from:
Russell and Norvig (AIMA Book): Chapter 6

- Part 1: Introduction
 - ① Introduction to AI ✓
 - ② Agents ✓
- Part 2: Problem-solving
 - ① Search 1: Uninformed Search ✓
 - ② Search 2: Heuristic Search ✓
 - ③ Search 3: Local Search ✓
 - ④ Search 4: Adversarial Search ✓
- Part 3: Reasoning and Uncertainty
 - ① **Reasoning 1: Constraint Satisfaction**
 - ② Reasoning 2: Logic and Inference
 - ③ Probabilistic Reasoning 1: BNs
 - ④ Probabilistic Reasoning 2: HMMs
- Part 4: Planning
 - ① Planning 1: Intro and Formalism
 - ② Planning 2: Algos and Heuristics
 - ③ Planning 3: Hierarchical Planning
 - ④ Planning 4: Stochastic Planning
- Part 5: Learning
 - ① Learning 1: Intro to ML
 - ② Learning 2: Regression
 - ③ Learning 3: Neural Networks
 - ④ Learning 4: Reinforcement Learning
- Part 6: Conclusion
 - ① Ethical Issues in AI
 - ② Conclusions and Discussion

- Defining Constraint Satisfaction Problems (CSP) ✓
- CSP examples ✓
- Backtracking search for CSPs ✓
- Local search for CSPs
- Problem structure and problem decomposition



Outline


1 Local Search for CSPs

► Local Search for CSPs

- Local search algorithms can be very effective in solving many CSPs.
- Local search algorithms use a complete-state formulation where:
 - each state assigns a value to every variable, and
 - the search changes the value of one variable at a time.
- **Min-conflicts heuristic:** value that results in the minimum number of conflicts with other variables that brings us closer to a solution.
 - Usually has a series of plateaus
- **Plateau search:** allowing sideways moves to another state with the same score.
 - can help local search find its way off the plateau.
- **Constraint weighting** aims to concentrate the search on the important constraints
 - Each constraint is given a numeric weight, initially all 1.
 - weights adjusted by incrementing when it is violated by the current assignment


```
1: function Min-Conflicts(csp,max_steps) returns a solution or failure
2:   current  $\leftarrow$  an initial complete assignment for csp
3:   for i = 1 to max_steps do
4:     if current is a solution for csp then return current
5:     var  $\leftarrow$  a randomly chosen conflicted variable from csp.Variables
6:     value  $\leftarrow \arg \min_{v \in var.Domain} \text{Conflicts}(csp, var, v, current)$ 
7:     current.var  $\leftarrow$  value
8:   return failure
```


- **States:** 4 queens in 4 columns ($4^4 = 256$ states)
- **Operators:** move queen in column
- **Goal test:** no attacks
- **Evaluation:** $h(n) = \text{number of attacks}$



- Given random initial state, can solve n -queens in almost constant time for arbitrary n with high probability (e.g., $n = 10,000,000$)
- The same appears to be true for any randomly-generated CSP **except** in a narrow range of the ratio:

$$R = \frac{\text{number of constraints}}{\text{number of variables}}$$






- Tasmania and mainland are **independent subproblems**
- Identifiable as **connected components** of constraint graph

- Suppose each subproblem has c variables out of n total
- Worst-case solution cost is $n/c \cdot d^c$, linear in n

E.g. $n = 80, d = 2, c = 20$


$2^{80} = 4$ billion years at 10 million nodes/sec

$4 \cdot 2^{20} = 0.4$ seconds at 10 million nodes/sec




- **Theorem:** if the constraint graph has no loops, the CSP can be solved in $O(nd^2)$ time
- Compare to general CSPs, where worst-case time is $O(d^n)$
- This property also applies to logical and probabilistic reasoning:
an important example of the relation between syntactic restrictions and the complexity of reasoning.

- 1 Choose a variable as root, order variables from root to leaves such that every node's parent precedes it in the ordering



- 2 For j from n down to 2, apply RemoveInconsistent($\text{Parent}(X_j), X_j$)
- 3 For j from 1 to n , assign X_j consistently with $\text{Parent}(X_j)$

Conditioning: instantiate a variable, prune its neighbours' domains



Cutset conditioning: instantiate (in all ways) a set of variables such that the remaining constraint graph is a tree

Cutset size $c \Rightarrow$ runtime $O(d^c \cdot (n - c)d^2)$, very fast for small c

- Hill-climbing, simulated annealing typically work with “complete” states, i.e., all variables assigned
- To apply to CSPs:
 - allow states with unsatisfied constraints
 - operators **reassign** variable values
- Variable selection: randomly select any conflicted variable
- Value selection by **min-conflicts** heuristic:
 - choose value that violates the fewest constraints
 - i.e., hillclimb with $h(n) = \text{total number of violated constraints}$

- CSPs are a special kind of problem:
 - states defined by values of a fixed set of variables
 - goal test defined by constraints on variable values
- Backtracking = depth-first search with one variable assigned per node
- Variable ordering and value selection heuristics help significantly
- Forward checking prevents assignments that guarantee later failure
- Constraint propagation (e.g., arc consistency) does additional work to constrain values and detect inconsistencies
- Local search using the min-conflicts heuristic has also been applied to constraint satisfaction problems with great success
- The CSP representation allows analysis of problem structure
- Tree-structured CSPs can be solved in linear time
- CSPs still **beat the most advanced ML methods** for most optimization tasks

- Modelling Constraint Satisfaction Problems
 - Solving CSP via Search
 - Heuristics
 - Solving CSP via Local Search



UNIVERSITY OF
ABERDEEN



Any Questions.