

Artificial Intelligence Foundation - **JC3001**

Lecture 23: Uncertainty over Time II

Prof. Aladdin Ayesh (aladdin.ayesh@abdn.ac.uk)

Dr. Binod Bhattacharai (binod.bhattacharai@abdn.ac.uk)

Dr. Gideon Ogunniye, (g.ogunniye@abdn.ac.uk)

September 2025



UNIVERSITY OF
ABERDEEN



Material adapted from:
Russell and Norvig (AIMA Book): Chapter 14 (14.1-14.3)

- Part 1: Introduction
 - ① Introduction to AI ✓
 - ② Agents ✓
- Part 2: Problem-solving
 - ① Search 1: Uninformed Search ✓
 - ② Search 2: Heuristic Search ✓
 - ③ Search 3: Local Search ✓
 - ④ Search 4: Adversarial Search ✓
- Part 3: Reasoning and Uncertainty
 - ① Reasoning 1: Constraint Satisfaction ✓
 - ② Reasoning 2: Logic and Inference ✓
 - ③ Probabilistic Reasoning 1: BNs ✓
 - ④ **Probabilistic Reasoning 2: HMMs**
- Part 4: Planning
 - ① Planning 1: Intro and Formalism
 - ② Planning 2: Algos and Heuristics
 - ③ Planning 3: Hierarchical Planning
 - ④ Planning 4: Stochastic Planning
- Part 5: Learning
 - ① Learning 1: Intro to ML
 - ② Learning 2: Regression
 - ③ Learning 3: Neural Networks
 - ④ Learning 4: Reinforcement Learning
- Part 6: Conclusion
 - ① Ethical Issues in AI
 - ② Conclusions and Discussion

- Time and Uncertainty ✓
- Inference in Temporal Models
- Hidden Markov Models




Outline

1 Inference in Temporal Models

► Inference in Temporal Models

► Hidden Markov Models



Smoothing computes $P(X_k | \mathbf{e}_{1:t})$, the posterior distribution of the state at some past time k given a complete sequence of observations from 1 to t .



UNIVERSITY OF
ABERDEEN

Inference in Temporal Models

Smoothing

$$P(X_k \mid \mathbf{e}_{1:t}) = P(X_k \mid \mathbf{e}_{1:k}, \mathbf{e}_{k+1:t})$$



UNIVERSITY OF
ABERDEEN

Inference in Temporal Models

Smoothing

$$\begin{aligned} P(X_k \mid \mathbf{e}_{1:t}) &= P(X_k \mid \mathbf{e}_{1:k}, \mathbf{e}_{k+1:t}) \\ &= \alpha P(X_k \mid \mathbf{e}_{1:k}) P(\mathbf{e}_{k+1:t} \mid X_k, \mathbf{e}_{1:k}) \end{aligned} \quad (\text{using Bayes' rule, given } \mathbf{e}_{1:k})$$



UNIVERSITY OF
ABERDEEN

Inference in Temporal Models

Smoothing

$$\begin{aligned} P(X_k \mid \mathbf{e}_{1:t}) &= P(X_k \mid \mathbf{e}_{1:k}, \mathbf{e}_{k+1:t}) \\ &= \alpha P(X_k \mid \mathbf{e}_{1:k}) P(\mathbf{e}_{k+1:t} \mid X_k, \mathbf{e}_{1:k}) && (\text{using Bayes' rule, given } \mathbf{e}_{1:k}) \\ &= \alpha P(X_k \mid \mathbf{e}_{1:k}) P(\mathbf{e}_{k+1:t} \mid X_k) && (\text{using conditional independence}) \end{aligned}$$



UNIVERSITY OF
ABERDEEN

Inference in Temporal Models

Smoothing

$$\begin{aligned} P(X_k \mid \mathbf{e}_{1:t}) &= P(X_k \mid \mathbf{e}_{1:k}, \mathbf{e}_{k+1:t}) \\ &= \alpha P(X_k \mid \mathbf{e}_{1:k}) P(\mathbf{e}_{k+1:t} \mid X_k, \mathbf{e}_{1:k}) && (\text{using Bayes' rule, given } \mathbf{e}_{1:k}) \\ &= \alpha P(X_k \mid \mathbf{e}_{1:k}) P(\mathbf{e}_{k+1:t} \mid X_k) && (\text{using conditional independence}) \\ &= \alpha \mathbf{f}_{1:k} \times \mathbf{b}_{k+1:t} \end{aligned}$$

- where “ \times ” represents pointwise multiplication of vectors.
- backward message $\mathbf{b}_{k+1:t}$ can be computed by recursive process that runs backward from t



UNIVERSITY OF
ABERDEEN

Inference in Temporal Models

Smoothing

$$\begin{aligned} P(X_k | \mathbf{e}_{1:t}) &= P(X_k | \mathbf{e}_{1:k}, \mathbf{e}_{k+1:t}) \\ &= \alpha P(X_k | \mathbf{e}_{1:k}) P(\mathbf{e}_{k+1:t} | X_k, \mathbf{e}_{1:k}) && (\text{using Bayes' rule, given } \mathbf{e}_{1:k}) \\ &= \alpha P(X_k | \mathbf{e}_{1:k}) P(\mathbf{e}_{k+1:t} | X_k) && (\text{using conditional independence}) \\ &= \alpha \mathbf{f}_{1:k} \times \mathbf{b}_{k+1:t} \end{aligned}$$

- where “ \times ” represents pointwise multiplication of vectors.
- backward message $\mathbf{b}_{k+1:t}$ can be computed by recursive process that runs backward from t

$$P(\mathbf{e}_{k+1:t} | X_k) = \sum_{x_{k+1}} P(\mathbf{e}_{k+1:t} | X_k, x_{k+1}) P(x_{k+1} | X_k) && (\text{conditioning on } X_{k+1})$$



UNIVERSITY OF
ABERDEEN

Inference in Temporal Models

Smoothing

$$\begin{aligned} P(X_k \mid \mathbf{e}_{1:t}) &= P(X_k \mid \mathbf{e}_{1:k}, \mathbf{e}_{k+1:t}) \\ &= \alpha P(X_k \mid \mathbf{e}_{1:k}) P(\mathbf{e}_{k+1:t} \mid X_k, \mathbf{e}_{1:k}) && \text{(using Bayes' rule, given } \mathbf{e}_{1:k}\text{)} \\ &= \alpha P(X_k \mid \mathbf{e}_{1:k}) P(\mathbf{e}_{k+1:t} \mid X_k) && \text{(using conditional independence)} \\ &= \alpha \mathbf{f}_{1:k} \times \mathbf{b}_{k+1:t} \end{aligned}$$

- where “ \times ” represents pointwise multiplication of vectors.
- backward message $\mathbf{b}_{k+1:t}$ can be computed by recursive process that runs backward from t

$$\begin{aligned} P(\mathbf{e}_{k+1:t} \mid X_k) &= \sum_{x_{k+1}} P(\mathbf{e}_{k+1:t} \mid X_k, x_{k+1}) P(x_{k+1} \mid X_k) && \text{(conditioning on } X_{k+1}\text{)} \\ &= \sum_{x_{k+1}} P(\mathbf{e}_{k+1:t} \mid x_{k+1}) P(x_{k+1} \mid X_k) && \text{(by conditional independence)} \end{aligned}$$

UNIVERSITY OF
ABERDEEN

Inference in Temporal Models

Smoothing

$$\begin{aligned} P(X_k \mid \mathbf{e}_{1:t}) &= P(X_k \mid \mathbf{e}_{1:k}, \mathbf{e}_{k+1:t}) \\ &= \alpha P(X_k \mid \mathbf{e}_{1:k}) P(\mathbf{e}_{k+1:t} \mid X_k, \mathbf{e}_{1:k}) && \text{(using Bayes' rule, given } \mathbf{e}_{1:k}\text{)} \\ &= \alpha P(X_k \mid \mathbf{e}_{1:k}) P(\mathbf{e}_{k+1:t} \mid X_k) && \text{(using conditional independence)} \\ &= \alpha \mathbf{f}_{1:k} \times \mathbf{b}_{k+1:t} \end{aligned}$$

- where “ \times ” represents pointwise multiplication of vectors.
- backward message $\mathbf{b}_{k+1:t}$ can be computed by recursive process that runs backward from t

$$\begin{aligned} P(\mathbf{e}_{k+1:t} \mid X_k) &= \sum_{x_{k+1}} P(\mathbf{e}_{k+1:t} \mid X_k, x_{k+1}) P(x_{k+1} \mid X_k) && \text{(conditioning on } X_{k+1}\text{)} \\ &= \sum_{x_{k+1}} P(\mathbf{e}_{k+1:t} \mid x_{k+1}) P(x_{k+1} \mid X_k) && \text{(by conditional independence)} \\ &= \sum_{x_{k+1}} P(\mathbf{e}_{k+1}, \mathbf{e}_{k+2:t} \mid x_{k+1}) P(x_{k+1} \mid X_k) \end{aligned}$$

UNIVERSITY OF
ABERDEEN

Inference in Temporal Models

Smoothing

$$\begin{aligned} P(X_k | \mathbf{e}_{1:t}) &= P(X_k | \mathbf{e}_{1:k}, \mathbf{e}_{k+1:t}) \\ &= \alpha P(X_k | \mathbf{e}_{1:k}) P(\mathbf{e}_{k+1:t} | X_k, \mathbf{e}_{1:k}) && \text{(using Bayes' rule, given } \mathbf{e}_{1:k}\text{)} \\ &= \alpha P(X_k | \mathbf{e}_{1:k}) P(\mathbf{e}_{k+1:t} | X_k) && \text{(using conditional independence)} \\ &= \alpha \mathbf{f}_{1:k} \times \mathbf{b}_{k+1:t} \end{aligned}$$

- where “ \times ” represents pointwise multiplication of vectors.
- backward message $\mathbf{b}_{k+1:t}$ can be computed by recursive process that runs backward from t

$$\begin{aligned} P(\mathbf{e}_{k+1:t} | X_k) &= \sum_{x_{k+1}} P(\mathbf{e}_{k+1:t} | X_k, x_{k+1}) P(x_{k+1} | X_k) && \text{(conditioning on } X_{k+1}\text{)} \\ &= \sum_{x_{k+1}} P(\mathbf{e}_{k+1:t} | x_{k+1}) P(x_{k+1} | X_k) && \text{(by conditional independence)} \\ &= \sum_{x_{k+1}} P(\mathbf{e}_{k+1}, \mathbf{e}_{k+2:t} | x_{k+1}) P(x_{k+1} | X_k) \\ &= \sum_{x_{k+1}} \underbrace{P(\mathbf{e}_{k+1} | x_{k+1})}_{\text{sensor model}} \underbrace{P(\mathbf{e}_{k+2:t} | x_{k+1})}_{\text{recursion}} \underbrace{P(x_{k+1} | X_k)}_{\text{transition model}} \end{aligned}$$

We now have a recursive formulation for the backward message:

$$P(\mathbf{e}_{k+1:t} \mid X_k) = \sum_{x_{k+1}} P(\mathbf{e}_{k+1} \mid x_{k+1}) P(\mathbf{e}_{k+2:t} \mid x_{k+1}) P(x_{k+1} \mid X_k)$$

Which we can now use as a function in our next algorithm:

$$\mathbf{b}_{k+1:t} = \text{Backward}(\mathbf{b}_{k+2:t}, \mathbf{e}_{k+1})$$

function Forward-Backward($ev, prior$) **returns** a vector of probability distributions

inputs: ev , a vector of evidence values for steps $1, \dots, t$

$prior$, the prior distribution on the initial state $P(X_0)$

local variables: fv , a vector of forward messages for steps $0, \dots, t$

b , a representation of the backward message, initially all 1s

sv , a vector of smoothed estimates for steps $1, \dots, t$

$fv[0] \leftarrow prior$

for $i = 1$ **to** t **do**

$fv[i] \leftarrow \text{Forward}(fv[i - 1], ev[i])$

for $i = t$ **down to** 1 **do**

$sv[i] \leftarrow \text{Normalize}(fv[i] \times b)$

$b \leftarrow \text{Backward}(b, ev[i])$

return sv

The forward-backward algorithm for smoothing: computing posterior probabilities of a sequence of states given a sequence of observations



Smoothing Example

1 Inference in Temporal Models

Back to the umbrella example, let us compute the smoothed estimate for rain on day one, given we saw an umbrella on days one and two, given by:

$$P(R_1 \mid u_1, u_2) = \alpha P(R_1 \mid u_1) P(u_2 \mid R_1)$$

Back to the umbrella example, let us compute the smoothed estimate for rain on day one, given we saw an umbrella on days one and two, given by:

$$P(R_1 | u_1, u_2) = \alpha P(R_1 | u_1) P(u_2 | R_1)$$

We know (from our filtering example), that $P(R_1 | u_1) = \langle 0.818, 0.182 \rangle$, so we need to compute the backward recursion

$$P(u_2 | R_1) = \sum_{r_2} P(u_2 | r_2) P(r_2 | R_1)$$

Back to the umbrella example, let us compute the smoothed estimate for rain on day one, given we saw an umbrella on days one and two, given by:

$$P(R_1 | u_1, u_2) = \alpha P(R_1 | u_1) P(u_2 | R_1)$$

We know (from our filtering example), that $P(R_1 | u_1) = \langle 0.818, 0.182 \rangle$, so we need to compute the backward recursion

$$\begin{aligned} P(u_2 | R_1) &= \sum_{r_2} P(u_2 | r_2) P(r_2 | R_1) \\ &= (0.9 \times 1 \times \langle 0.7, 0.3 \rangle) + (0.2 \times 1 \times \langle 0.3, 0.7 \rangle) \end{aligned}$$

Back to the umbrella example, let us compute the smoothed estimate for rain on day one, given we saw an umbrella on days one and two, given by:

$$P(R_1 | u_1, u_2) = \alpha P(R_1 | u_1) P(u_2 | R_1)$$

We know (from our filtering example), that $P(R_1 | u_1) = \langle 0.818, 0.182 \rangle$, so we need to compute the backward recursion

$$\begin{aligned} P(u_2 | R_1) &= \sum_{r_2} P(u_2 | r_2) P(r_2 | R_1) \\ &= (0.9 \times 1 \times \langle 0.7, 0.3 \rangle) + (0.2 \times 1 \times \langle 0.3, 0.7 \rangle) = \langle 0.69, 0.41 \rangle \end{aligned}$$

Back to the umbrella example, let us compute the smoothed estimate for rain on day one, given we saw an umbrella on days one and two, given by:

$$P(R_1 | u_1, u_2) = \alpha P(R_1 | u_1) P(u_2 | R_1)$$

We know (from our filtering example), that $P(R_1 | u_1) = \langle 0.818, 0.182 \rangle$, so we need to compute the backward recursion

$$\begin{aligned} P(u_2 | R_1) &= \sum_{r_2} P(u_2 | r_2) P(r_2 | R_1) \\ &= (0.9 \times 1 \times \langle 0.7, 0.3 \rangle) + (0.2 \times 1 \times \langle 0.3, 0.7 \rangle) = \langle 0.69, 0.41 \rangle \end{aligned}$$

Plugging this value in the equation above we have:

$$P(R_1 | u_1, u_2) = \alpha \langle 0.818, 0.182 \rangle \times \langle 0.69, 0.41 \rangle \approx \langle 0.883, 0.117 \rangle$$

Finding the most likely sequence

- There is a linear-time algorithm for finding the most likely sequence
- It relies on the same Markov property that yielded efficient algorithms for filtering and smoothing
- View each sequence as a path through a graph whose nodes are the possible states at each time step.
- likelihood of any path is the product of the transition probabilities along the path and the probabilities of the given observations at each state
- There is a recursive relationship between most likely paths to each state x_{t+1} and most likely paths to each state x_t



UNIVERSITY OF
ABERDEEN

Finding the most likely sequence

- Recursively computed message $m_{1:t}$

$$m_{1:t} = \max_{x_{1:t-1}}$$


$$\begin{aligned} m_{1:t+1} &= \max_{x_{1:t}} P(x_{1:t}, X_{t+1}, \mathbf{e}_{1:t+1}) = \max_{x_{1:t}} P(x_{1:t}, X_{t+1}, \mathbf{e}_{1:t}, \mathbf{e}_{t+1}) \\ &= \max_{x_{1:t}} P(\mathbf{e}_{t+1} \mid x_{t:1}, X_{t+1}, \mathbf{e}_{1:t}) P(x_{1:t}, X_{t+1}, \mathbf{e}_{1:t}) \\ &= P(\mathbf{e}_{t+1} \mid X_{t+1}) \max_{x_{1:t}} P(X_{t+1}, \mid x_t) P(x_{1:t}, \mathbf{e}_{1:t}) \\ &= P(\mathbf{e}_{t+1} \mid X_{t+1}) \max_{x_t} P(X_{t+1}, \mid x_t) \max_{x_{1:t-1}} P(x_{1:t-1}, x_t, \mathbf{e}_{1:t}) \end{aligned}$$

- $m_{1:t}$ will contain the probability for the most likely sequence reaching each of the final states.

Viterbi algorithm

- Select the final state of the most likely sequence overall.
- In order to identify the actual sequence, as opposed to just computing its probability, the algorithm will also need to record, for each state, the best state that leads to it.

1 Inference in Temporal Models



- (a) Possible state sequences for Rain can be viewed as paths through a graph of the possible states at each time step.
- (b) Operation of the Viterbi algorithm for the umbrella observation sequence [true; true; false; true; true], where the evidence starts at time 1.



Outline

2 Hidden Markov Models

- ▶ Inference in Temporal Models
- ▶ Hidden Markov Models

- Hidden Markov model, or HMM:

- temporal probabilistic model
- the state of the process is described by a single, discrete random variable
- No restriction on the evidence variables.

There can be many evidence variables, both discrete and continuous.

Simplified matrix algorithms

The transition model $P(X_t | X_{t-1})$ becomes an $S \times S$ matrix T where:

$$T_{i,j} = P(X_t = j | X_{t-1} = i)$$

$T_{i,j}$ is the probability of a transition from state i to state j

- Matrix formulation allows improved algorithms
 - simple variation on the forward-backward algorithm that allows smoothing to be carried out in constant space, independently of the length of the sequence
 - Online smoothing with a fixed lag.



function Fixed-lag-Smoothing(\mathbf{e} , hmm , d)

 Inputs: \mathbf{e}_t , the current evidence for time step t

hmm , a hidden Markov model with $S \times S$ transition matrix T

d , the length of the lag for smoothing

 persistent: t , the current time, initially 1

\mathbf{f} , the forward message $P(X_t | \mathbf{e}_{1:t})$, initially $hmm.Prior$

B , the d -step backward transformation matrix, initially the identity matrix

$\mathbf{e}_{t-d:t}$, double-ended list of evidence from $t - d$ to t , initially empty

 local variables: O_{t-d} , O_t , diagonal matrices containing the sensor model information

 add \mathbf{e}_t to the end of $\mathbf{e}_{t-d:t}$

$O_t \leftarrow$ diagonal matrix containing $P(\mathbf{e}_t | X_t)$

 if $t > d$ then

$\mathbf{f} \leftarrow \text{Forward}(\mathbf{f}, \mathbf{e}_{t-d})$

 remove \mathbf{e}_{t-d-1} from the beginning of $\mathbf{e}_{t-d:t}$

$O_{t-d} \leftarrow$ diagonal matrix containing $P(\mathbf{e}_{t-d} | X_{t-d})$

$B \leftarrow O_{t-d}^{-1} T^{-1} B T O_t$

 else $B \leftarrow B T O_t$

$t \leftarrow t + 1$

 if $t > d + 1$ then return Normalize($\mathbf{f} \times B$) else return null



Hidden Markov Models

2 Hidden Markov Models

An algorithm for smoothing with a fixed time lag of d steps, implemented as an online algorithm that outputs the new smoothed estimate given the observation for a new time step.

Robot Localization:


- Robot can start at any valid location of the 4×16 grid
- Moves randomly over the map following
- 42 Valid positions
- Sensor E_t has 16 values: 4-bit sequences of whether there is an obstacle: NESW

(a) Posterior distribution over robot location after $E_1 = 1011$ (b) Posterior distribution over robot location after $E_1 = 1011, E_2 = 1010$




Hidden Markov Models

Performance



(a)



(b)

HMM localization performance over the length of the observation sequence for various different sensor error probabilities ϵ ; data averaged over 400 runs.

- The localization error, defined as the Manhattan distance from the true location.
- The Viterbi path error, defined as the average Manhattan distance of states on the Viterbi path from corresponding states on the true path

- Temporal Models
 - Filtering
 - Prediction
 - Smoothing
 - Most likely explanation
- Viterbi's Algorithm
- Hidden Markov Models



UNIVERSITY OF
ABERDEEN



Any Questions.