



UNIVERSITY OF
ABERDEEN

Artificial Intelligence Foundation - JC3001

Lecture 44: Reinforcement Learning - I

Prof. Aladdin Ayesh (aladdin.ayesh@abdn.ac.uk)

Dr. Binod Bhattacharai (binod.bhattacharai@abdn.ac.uk)

Dr. Gideon Ogunniye, (g.ogunniye@abdn.ac.uk)

October 2025



Material adapted from:

Russell and Norvig (AIMA Book): Chapter 22

Sutton and Barto (Reinforcement Learning: An Introduction 2nd ed.)

David Silver (UCL)

Michael Littman (Brown University) and Charles Isbell (GA Tech)

Course Progression

- Part 1: Introduction
 - ① Introduction to AI ✓
 - ② Agents ✓
- Part 2: Problem-solving
 - ① Search 1: Uninformed Search ✓
 - ② Search 2: Heuristic Search ✓
 - ③ Search 3: Local Search ✓
 - ④ Search 4: Adversarial Search ✓
- Part 3: Reasoning and Uncertainty
 - ① Reasoning 1: Constraint Satisfaction ✓
 - ② Reasoning 2: Logic and Inference ✓
 - ③ Probabilistic Reasoning 1: BNs ✓
 - ④ Probabilistic Reasoning 2: HMMs ✓
- Part 4: Planning
 - ① Planning 1: Intro and Formalism ✓
 - ② Planning 2: Algorithms & Heuristics ✓
 - ③ Planning 3: Hierarchical Planning ✓
 - ④ Planning 4: Stochastic Planning ✓
- Part 5: Learning
 - ① Learning 1: Intro to ML ✓
 - ② Learning 2: Regression ✓
 - ③ Learning 3: Neural Networks ✓
 - ④ Learning 4: Reinforcement Learning
- Part 6: Conclusion
 - ① Ethical Issues in AI
 - ② Conclusions and Discussion



Objectives

- Bandit Problems
- Reinforcement Learning based Agents
- Tabular Reinforcement Learning
- Function Generalization



Outline

1 Recap and Motivation

- ▶ Recap and Motivation
- ▶ Multi-armed Bandits
- ▶ Reinforcement Learning based Agents

Motivation

1 Recap and Motivation

	1	2	3	4
a				+1
b				-1
c				

- Recall the grid world example we saw for MDPs

Motivation

1 Recap and Motivation

	1	2	3	4
a				+1
b				-1
c				

- Recall the grid world example we saw for MDPs
- We learned how to calculate the optimal policy using transition **probabilities** and **rewards**

Motivation

1 Recap and Motivation

	1	2	3	4
a				+1
b				-1
c				

- Recall the grid world example we saw for MDPs
- We learned how to calculate the optimal policy using transition **probabilities** and **rewards**
- But what if we do not know the rewards, probabilities or both?

Motivation

1 Recap and Motivation

	1	2	3	4
a				+1
b				-1
c				

- Recall the grid world example we saw for MDPs
- We learned how to calculate the optimal policy using transition **probabilities** and **rewards**
- But what if we do not know the rewards, probabilities or both?
- We should be able to explore the environment and discover the best policy

Learning Recap

1 Recap and Motivation

Main methods of learning:

- Supervised learning:

From $(\vec{x}^{(1)}, y^{(1)}), (\vec{x}^{(2)}, y^{(2)}), \dots (\vec{x}^{(m)}, y^{(m)})$, derive a function $f(X^{(m)}) = Y^{(m)}$

Main methods of learning:

- Supervised learning:

From $(\vec{x}^{(1)}, y^{(1)}), (\vec{x}^{(2)}, y^{(2)}), \dots, (\vec{x}^{(m)}, y^{(m)})$, derive a function $f(X^{(m)}) = Y^{(m)}$

- Unsupervised learning:

From $\vec{x}^{(1)}, \vec{x}^{(2)}, \vec{x}^{(m)}, \dots, \vec{x}^{(m)}$, identify classes x_c and
derive a function to calculate $P(X = x_c)$

Learning Recap

1 Recap and Motivation

Main methods of learning:

- Supervised learning:

From $(\vec{x}^{(1)}, y^{(1)}), (\vec{x}^{(2)}, y^{(2)}), \dots, (\vec{x}^{(m)}, y^{(m)})$, derive a function $f(X^{(m)}) = Y^{(m)}$

- Unsupervised learning:

From $\vec{x}^{(1)}, \vec{x}^{(2)}, \vec{x}^{(m)}, \dots, \vec{x}^{(m)}$, identify classes x_c and
derive a function to calculate $P(X = x_c)$

- Reinforcement learning:

From a sequence of state transitions (i.e. states and actions)
and occasional rewards $[s \xrightarrow{r}, a, s \xrightarrow{r}, a, \dots]$,
discover the optimal policy $\pi^*(s)$ for this environment.

Machine Learning Problems

1 Recap and Motivation

Which of the previous methods of learning would be more suitable to solve the following problems:

Sup. Unsup. Reinf.

- Spam Filtering
- Speech Recognition
- Starlight classification
- Playing SF Ξ against an unknown player
- Navigating an unknown environment
- Elevator controller

Machine Learning Problems

1 Recap and Motivation

Which of the previous methods of learning would be more suitable to solve the following problems:

Sup. Unsup. Reinf.

X

- Spam Filtering
- Speech Recognition
- Starlight classification
- Playing SF Ξ against an unknown player
- Navigating an unknown environment
- Elevator controller

Machine Learning Problems

1 Recap and Motivation

Which of the previous methods of learning would be more suitable to solve the following problems:

Sup. Unsup. Reinf.

X

Spam Filtering

X

Speech Recognition

Starlight classification

Playing SF Ξ against an unknown player

Navigating an unknown environment

Elevator controller

Machine Learning Problems

1 Recap and Motivation

Which of the previous methods of learning would be more suitable to solve the following problems:

Sup. Unsup. Reinf.

X

Spam Filtering

X

Speech Recognition

X

Starlight classification

Playing SF \exists against an unknown player

Navigating an unknown environment

Elevator controller

Machine Learning Problems

1 Recap and Motivation

Which of the previous methods of learning would be more suitable to solve the following problems:

Sup. Unsup. Reinf.

X

Spam Filtering

X

Speech Recognition

X

Starlight classification

X

Playing SF \exists against an unknown player

Navigating an unknown environment

Elevator controller

Machine Learning Problems

1 Recap and Motivation

Which of the previous methods of learning would be more suitable to solve the following problems:

Sup. Unsup. Reinf.

X

Spam Filtering

X

Speech Recognition

X

Starlight classification

X

Playing SF \exists against an unknown player

X

Navigating an unknown environment

Elevator controller

Machine Learning Problems

1 Recap and Motivation

Which of the previous methods of learning would be more suitable to solve the following problems:

Sup. Unsup. Reinf.

X

Spam Filtering

X

Speech Recognition

X

Starlight classification

X

Playing SF \exists against an unknown player

X


Navigating an unknown environment

X

Elevator controller

Markov Decision Process (MDP)


1 Recap and Motivation



- An MDP is defined in terms of
 - ① An initial state S_0
 - ② A transition model $T(s, a, s') = P(s'|a, s)$ (Markovian)
 - ③ A reward function $R(s, a, s')$
- A solution to a MDP must specify what the agent should do for any state. Such a solution is called a policy

Markov Decision Process (MDP)

1 Recap and Motivation



- An MDP is defined in terms of
 - ① An initial state S_0
 - ② A transition model $T(s, a, s') = P(s'|a, s)$ (Markovian)
 - ③ A reward function $R(s, a, s')$ — sometimes expressed as $R(s')$
- A solution to a MDP must specify what the agent should do for any state. Such a solution is called a policy

Finding the Optimal Policy

1 Recap and Motivation

$$V(s) \leftarrow \left[\max_a \gamma \sum_{s'} P(s'|s, a) * V(s') \right] + R(s)$$

Finding the Optimal Policy

1 Recap and Motivation

$$V(s) \leftarrow \left[\max_a \gamma \sum_{s'} P(s'|s, a) * V(s') \right] + R(s)$$

$$\pi^*(s) = \arg \max_a \sum_{s'} P(s'|s, a) * V(s')$$



Outline

2 Multi-armed Bandits

- ▶ Recap and Motivation
- ▶ Multi-armed Bandits
- ▶ Reinforcement Learning based Agents

k-Armed Bandit Problem

2 Multi-armed Bandits

k-armed Bandit Problem

- Repeated choices among k different options or actions
- Each choice has a reward drawn from a stationary probability distribution
- Goal is to maximize rewards over a period of time



Thus, the value is the expected reward:

$$\begin{aligned} q_*(a) &\doteq \mathbb{E}[R_t \mid A_t = a] \quad \forall a \in \{1, \dots, k\} \\ &= \sum_r p(r \mid a)r \end{aligned}$$

And to maximize it over time

$$\arg \max_a q_*(a)$$

Computing the expected reward

2 Multi-armed Bandits

Consider a doctor with 3 possible treatments to offer a patient.

- Each treatment can either cure the patient ($r = 1$), do nothing ($r = 0$), or make the patient sicker ($r = -1$), with probabilities ($p(r | \text{drug})$) as follows:

d_1 🧑‍⚕️ $p(r = 1 | d = d_1) = 0.1$, $p(r = 0 | d = d_1) = 0.8$, and $p(r = -1 | d = d_1) = 0.1$

d_2 🧑‍⚕️ $p(r = 1 | d = d_2) = 0.3$, $p(r = 0 | d = d_2) = 0.1$, and $p(r = -1 | d = d_2) = 0.6$

d_3 🧑‍⚕️ $p(r = 1 | d = d_3) = 0.3$, $p(r = 0 | d = d_3) = 0.5$, and $p(r = -1 | d = d_3) = 0.2$

Computing the expected reward

2 Multi-armed Bandits

Consider a doctor with 3 possible treatments to offer a patient.

- Each treatment can either cure the patient ($r = 1$), do nothing ($r = 0$), or make the patient sicker ($r = -1$), with probabilities ($p(r | \text{drug})$) as follows:

$$d_1 \quad \text{☒ } p(r = 1 | d = d_1) = 0.1, p(r = 0 | d = d_1) = 0.8, \text{ and } p(r = -1 | d = d_1) = 0.1$$

$$d_2 \quad \text{☒ } p(r = 1 | d = d_2) = 0.3, p(r = 0 | d = d_2) = 0.1, \text{ and } p(r = -1 | d = d_2) = 0.6$$

$$d_3 \quad \text{☒ } p(r = 1 | d = d_3) = 0.3, p(r = 0 | d = d_3) = 0.5, \text{ and } p(r = -1 | d = d_3) = 0.2$$

Compute q_* for all drug options and help the doctor choose:

$$d_1 \quad \text{☒ } q(d = d_1) = (0.1 * 1) + (0.8 * 0) + (0.1 * -1) = 0$$

$$d_2 \quad \text{☒ } q(d = d_2) = (0.3 * 1) + (0.1 * 0) + (0.6 * -1) = -0.3$$

$$d_3 \quad \text{☒ } q(d = d_3) = (0.3 * 1) + (0.5 * 0) + (0.2 * -1) = 0.1$$

Sample-Average Method

$$Q_t(a) \doteq \frac{\text{sum of rewards for } a \text{ taken prior to } t}{\text{number of times } a \text{ taken prior to } t}$$
$$= \frac{\sum_{i=1}^{t-1} R_i}{t - 1}$$

Example

2 Multi-armed Bandits

$$\frac{\sum_{i=1}^{t-1} R_i}{t - 1}$$



$$Q_0(\text{blue}) = 0$$



$$Q_0(\text{red}) = 0$$



$$Q_0(\text{green}) = 0$$

Example

2 Multi-armed Bandits

$$\frac{\sum_{i=1}^{t-1} R_i}{t - 1}$$



1
 $Q_1(\text{medicine}) = 1$



o
 $Q_1(\text{medicine}) = 0$



o
 $Q_1(\text{medicine}) = 0$



Example

2 Multi-armed Bandits

$$\frac{\sum_{i=1}^{t-1} R_i}{t - 1}$$



1

o

$$Q_2(\text{medicine}) = 0.5$$



o

o

$$Q_2(\text{medicine}) = 0$$



o

o

$$Q_2(\text{medicine}) = 0$$



Example

2 Multi-armed Bandits

$$\frac{\sum_{i=1}^{t-1} R_i}{t - 1}$$



1

0

0

$$Q_3(\text{medicine}) = 0.5$$



0

0

-1

$$Q_3(\text{red liquid}) = -1$$



0

0

0

$$Q_3(\text{green liquid}) = 0$$



Example

2 Multi-armed Bandits



$$\frac{\sum_{i=1}^{t-1} R_i}{t-1}$$

1

0

0

0

0

0

0

-1

0

$$Q_4(\text{□}) = 0.5$$

$$Q_4(\text{□} \textcolor{red}{□}) = -1$$

$$Q_4(\text{□} \textcolor{green}{□}) = 1$$



Example

2 Multi-armed Bandits



$$\frac{\sum_{i=1}^{t-1} R_i}{t - 1}$$

1

0

0

0

0

0

0

-1

0

0

0

0

0

1

0

$$Q_5(\text{pill}) = 0.5$$

$$Q_5(\text{pill}) = -0.5$$

$$= Q_5(\text{pill}) = 1$$



Incremental Update Rule

2 Multi-armed Bandits

$$Q_{n+1} = \frac{1}{n} \sum_{i=1}^n R_i$$

Incremental Update Rule

2 Multi-armed Bandits

$$\begin{aligned} Q_{n+1} &= \frac{1}{n} \sum_{i=1}^n R_i \\ &= \frac{1}{n} (R_n + \sum_{i=1}^{n-1} R_i) \end{aligned}$$

Incremental Update Rule

2 Multi-armed Bandits

$$\begin{aligned} Q_{n+1} &= \frac{1}{n} \sum_{i=1}^n R_i \\ &= \frac{1}{n} (R_n + \sum_{i=1}^{n-1} R_i) \\ &= \frac{1}{n} (R_n + (n-1) \frac{1}{n-1} \sum_{i=1}^{n-1} R_i) \end{aligned}$$

Incremental Update Rule

2 Multi-armed Bandits

$$\begin{aligned}
 Q_{n+1} &= \frac{1}{n} \sum_{i=1}^n R_i \\
 &= \frac{1}{n} (R_n + \sum_{i=1}^{n-1} R_i) \\
 &= \frac{1}{n} (R_n + (n-1) \frac{1}{n-1} \sum_{i=1}^{n-1} R_i)
 \end{aligned}$$

$$\text{Recall: } Q_n = \frac{1}{n-1} \sum_{i=1}^{n-1} R_i$$

Incremental Update Rule

2 Multi-armed Bandits

$$\begin{aligned}
 Q_{n+1} &= \frac{1}{n} \sum_{i=1}^n R_i \\
 &= \frac{1}{n} (R_n + \sum_{i=1}^{n-1} R_i) \\
 &= \frac{1}{n} (R_n + (n-1) \frac{1}{n-1} \sum_{i=1}^{n-1} R_i)
 \end{aligned}$$

$$\begin{aligned}
 &= \frac{1}{n} (R_n + (n-1)Q_n) \\
 &= \frac{1}{n} (R_n + nQ_n - Q_n)
 \end{aligned}$$

$$\text{Recall: } Q_n = \frac{1}{n-1} \sum_{i=1}^{n-1} R_i$$

Incremental Update Rule

2 Multi-armed Bandits

$$\begin{aligned}
 Q_{n+1} &= \frac{1}{n} \sum_{i=1}^n R_i \\
 &= \frac{1}{n} (R_n + \sum_{i=1}^{n-1} R_i) \\
 &= \frac{1}{n} (R_n + (n-1) \frac{1}{n-1} \sum_{i=1}^{n-1} R_i) \\
 &= \frac{1}{n} (R_n + (n-1)Q_n) \\
 &= \frac{1}{n} (R_n + nQ_n - Q_n) \\
 &= Q_n + \frac{1}{n} (R_n - Q_n)
 \end{aligned}$$

$$\text{Recall: } Q_n = \frac{1}{n-1} \sum_{i=1}^{n-1} R_i$$

Incremental Update Rule

2 Multi-armed Bandits

We will see this form of update throughout this lecture:

$$\text{NewEstimate} \leftarrow \text{OldEstimate} + \text{StepSize} [\text{Target} - \text{OldEstimate}]$$

$$Q_{n+1} = Q_n + \frac{1}{n}(R_n - Q_n)$$

$$Q_{n+1} = Q_n + \alpha_n(R_n - Q_n)$$



Outline

3 Reinforcement Learning based Agents

- ▶ Recap and Motivation
- ▶ Multi-armed Bandits
- ▶ Reinforcement Learning based Agents

Reinforcement Learning based Agents

3 Reinforcement Learning based Agents

Agents execute **trials** in the environment, using some policy π .

Typical trials look like:

$$(1, 1)_{-.04} \rightsquigarrow (1, 2)_{-.04} \rightsquigarrow (1, 3)_{-.04} \rightsquigarrow (1, 2)_{-.04} \rightsquigarrow (1, 3)_{-.04} \rightsquigarrow (2, 3)_{-.04} \rightsquigarrow \\ (3, 3)_{-.04} \rightsquigarrow (4, 3)_{+1}$$

$$(1, 1)_{-.04} \rightsquigarrow (1, 2)_{-.04} \rightsquigarrow (1, 3)_{-.04} \rightsquigarrow (2, 3)_{-.04} \rightsquigarrow (3, 3)_{-.04} \rightsquigarrow (3, 2)_{-.04} \rightsquigarrow \\ (3, 3)_{-.04} \rightsquigarrow (4, 3)_{+1}$$

$$(1, 1)_{-.04} \rightsquigarrow (2, 1)_{-.04} \rightsquigarrow (3, 1)_{-.04} \rightsquigarrow (3, 2)_{-.04} \rightsquigarrow (4, 2)_{-1}$$



Agent Types for RL

3 Reinforcement Learning based Agents

What to do if we do not know R



Agent Types for RL

3 Reinforcement Learning based Agents

What to do if we do not know R , or P

Agent Types for RL

3 Reinforcement Learning based Agents

What to do if we do not know R , or P , or both?

- We can try to learn R and P

Agent Types for RL

3 Reinforcement Learning based Agents

What to do if we do not know R , or P , or both?

- We can try to learn R and P
- Or something else that obviates the need to learn them

Agent Types for RL

3 Reinforcement Learning based Agents

What to do if we do not know R , or P , or both?

- We can try to learn R and P
- Or something else that obviates the need to learn them

Agent Types for RL

3 Reinforcement Learning based Agents

What to do if we do not know R , or P , or both?

- We can try to learn R and P
- Or something else that obviates the need to learn them

Agent	We Know	What to Learn	What to use
-------	---------	---------------	-------------

Agent Types for RL

3 Reinforcement Learning based Agents

What to do if we do not know R , or P , or both?

- We can try to learn R and P
- Or something else that obviates the need to learn them

Agent	We Know	What to Learn	What to use
Utility-based agent	P	$R \rightarrow U$	U

Agent Types for RL

3 Reinforcement Learning based Agents

What to do if we do not know R , or P , or both?

- We can try to learn R and P
- Or something else that obviates the need to learn them

Agent	We Know	What to Learn	What to use
Utility-based agent	P	$R \rightarrow U$	U
Q-learning agent		$Q(s, a)$	Q

Agent Types for RL

3 Reinforcement Learning based Agents

What to do if we do not know R , or P , or both?

- We can try to learn R and P
- Or something else that obviates the need to learn them

Agent	We Know	What to Learn	What to use
Utility-based agent	P	$R \rightarrow U$	U
Q-learning agent		$Q(s, a)$	Q
Reflex agent		$\pi(s)$	π

Passive vs. Active

3 Reinforcement Learning based Agents

Agent	We Know	What to Learn	What to use
Utility-based agent	P	$R \rightarrow U$	U
Q-learning agent		$Q(s, a)$	Q
Reflex agent		$\pi(s)$	π

- Besides these designs, we decide how “adventurous” our agent will be

Passive vs. Active

3 Reinforcement Learning based Agents

Agent	We Know	What to Learn	What to use
Utility-based agent	P	$R \rightarrow U$	U
Q-learning agent		$Q(s, a)$	Q
Reflex agent		$\pi(s)$	π

- Besides these designs, we decide how “adventurous” our agent will be

Passive RL The agent simply follows a fixed policy, and learns the components above as he goes
The goal is to learn how good that policy is

Passive vs. Active

3 Reinforcement Learning based Agents

Agent	We Know	What to Learn	What to use
Utility-based agent	P	$R \rightarrow U$	U
Q-learning agent		$Q(s, a)$	Q
Reflex agent		$\pi(s)$	π

- Besides these designs, we decide how “adventurous” our agent will be

Passive RL The agent simply follows a fixed policy, and learns the components above as he goes
 The goal is to learn how good that policy is

Active RL After learning the components for a while,
 the agent changes policy to “cash in” on the learned components



UNIVERSITY OF
ABERDEEN



To continue in the next session.