



Artificial Intelligence Foundation - JC3001

Lecture 45: Reinforcement Learning - II

Prof. Aladdin Ayesh (aladdin.ayesh@abdn.ac.uk)

Dr. Binod Bhattacharai (binod.bhattacharai@abdn.ac.uk)

Dr. Gideon Ogunniye, (g.ogunniye@abdn.ac.uk)

October 2025



Material adapted from:
Russell and Norvig (AIMA Book): Chapter 22
Sutton and Barto (Reinforcement Learning: An Introduction 2nd ed.)
David Silver (UCL)
Michael Littman (Brown University) and Charles Isbell (GA Tech)

Course Progression

- Part 1: Introduction
 - ① Introduction to AI ✓
 - ② Agents ✓
- Part 2: Problem-solving
 - ① Search 1: Uninformed Search ✓
 - ② Search 2: Heuristic Search ✓
 - ③ Search 3: Local Search ✓
 - ④ Search 4: Adversarial Search ✓
- Part 3: Reasoning and Uncertainty
 - ① Reasoning 1: Constraint Satisfaction ✓
 - ② Reasoning 2: Logic and Inference ✓
 - ③ Probabilistic Reasoning 1: BNs ✓
 - ④ Probabilistic Reasoning 2: HMMs ✓
- Part 4: Planning
 - ① Planning 1: Intro and Formalism ✓
 - ② Planning 2: Algos and Heuristics ✓
 - ③ Planning 3: Hierarchical Planning ✓
 - ④ Planning 4: Stochastic Planning ✓
- Part 5: Learning
 - ① Learning 1: Intro to ML ✓
 - ② Learning 2: Regression ✓
 - ③ Learning 3: Neural Networks ✓
 - ④ Learning 4: Reinforcement Learning
- Part 6: Conclusion
 - ① Ethical Issues in AI
 - ② Conclusions and Discussion

Objectives

- Bandit Problems ✓
- Reinforcement Learning based Agents ✓
- Tabular Reinforcement Learning
- Function Generalization



Outline

1 TD Learning

- ▶ TD Learning
- ▶ Q Learning
- ▶ Feature Generalization

Passive Temporal Difference Learning

1 TD Learning

- We start with a policy π

	1	2	3	4
a				+1
b				-1
c				

Algorithm consists of:

if s' is new **then** $U[s'] \leftarrow r'$

if s is not null **then**

increment $N_s[s]$

$U[s] \leftarrow U[s] + \alpha(N_s[s])(r + \gamma U[s'] - U[s])$

Passive Temporal Difference Learning

1 TD Learning

- We start with a policy π
- $U(s)$ — Utility for s (init. null)

	1	2	3	4
a				+1
b				-1
c				

Algorithm consists of:

if s' is new **then** $U[s'] \leftarrow r'$

if s is not null **then**

increment $N_s[s]$

$U[s] \leftarrow U[s] + \alpha(N_s[s])(r + \gamma U[s'] - U[s])$

Passive Temporal Difference Learning

1 TD Learning

- We start with a policy π
- $U(s)$ — Utility for s (init. null)
- $N(s)$ — # times we visited s

	1	2	3	4
a				+1
b				-1
c				

Algorithm consists of:

if s' is new **then** $U[s'] \leftarrow r'$

if s is not null **then**

increment $N_s[s]$

$U[s] \leftarrow U[s] + \alpha(N_s[s])(r + \gamma U[s'] - U[s])$

Passive Temporal Difference Learning

- We start with a policy π
- $U(s)$ — Utility for s (init. null)
- $N(s)$ — # times we visited s

Algorithm consists of:

- Follow π from the initial state to an end state (multiple times);

1 TD Learning

	1	2	3	4
a				+1
b				-1
c				

if s' is new **then** $U[s'] \leftarrow r'$

if s is not null **then**

increment $N_s[s]$

$U[s] \leftarrow U[s] + \alpha(N_s[s])(r + \gamma U[s'] - U[s])$

Passive Temporal Difference Learning

1 TD Learning

- We start with a policy π
- $U(s)$ — Utility for s (init. null)
- $N(s)$ — # times we visited s

	1	2	3	4
a				+1
b				-1
c				

Algorithm consists of:

- Follow π from the initial state to an end state (multiple times);
- Keep track of the number of times we visited each state;

if s' is new **then** $U[s'] \leftarrow r'$

if s is not null **then**

increment $N_s[s]$

$U[s] \leftarrow U[s] + \alpha(N_s[s])(r + \gamma U[s'] - U[s])$

Passive Temporal Difference Learning

1 TD Learning

- We start with a policy π
- $U(s)$ — Utility for s (init. null)
- $N(s)$ — # times we visited s

	1	2	3	4
a				+1
b				-1
c				

Algorithm consists of:

- Follow π from the initial state to an end state (multiple times);
- Keep track of the number of times we visited each state;
- Update the utility of the states as we go, using the following algorithm:

if s' is new **then** $U[s'] \leftarrow r'$

if s is not null **then**

increment $N_s[s]$

$U[s] \leftarrow U[s] + \alpha(N_s[s])(r + \gamma U[s'] - U[s])$

Passive Temporal Difference Learning

1 TD Learning

- We start with a policy π
- $U(s)$ — Utility for s (init. null)
- $N(s)$ — # times we visited s

	1	2	3	4
a				+1
b	0			-1
c				

Algorithm consists of:

- Follow π from the initial state to an end state (multiple times);
- Keep track of the number of times we visited each state;
- Update the utility of the states as we go, using the following algorithm:

if s' is new **then** $U[s'] \leftarrow r'$

if s is not null **then**

increment $N_s[s]$

$U[s] \leftarrow U[s] + \alpha(N_s[s])(r + \gamma U[s'] - U[s])$

Passive Temporal Difference Learning

1 TD Learning

- We start with a policy π
- $U(s)$ — Utility for s (init. null)
- $N(s)$ — # times we visited s

	1	2	3	4
a	0			+1
b	0			-1
c				

Algorithm consists of:

- Follow π from the initial state to an end state (multiple times);
- Keep track of the number of times we visited each state;
- Update the utility of the states as we go, using the following algorithm:
if s' is new **then** $U[s'] \leftarrow r'$
if s is not null **then**

increment $N_s[s]$

$$U[s] \leftarrow U[s] + \alpha(N_s[s])(r + \gamma U[s'] - U[s])$$

Passive Temporal Difference Learning

1 TD Learning

- We start with a policy π
- $U(s)$ — Utility for s (init. null)
- $N(s)$ — # times we visited s

	1	2	3	4
a	0	0		+1
b	0			-1
c				

Algorithm consists of:

- Follow π from the initial state to an end state (multiple times);
- Keep track of the number of times we visited each state;
- Update the utility of the states as we go, using the following algorithm:
if s' is new **then** $U[s'] \leftarrow r'$
if s is not null **then**
 increment $N_s[s]$

$$U[s] \leftarrow U[s] + \alpha(N_s[s])(r + \gamma U[s'] - U[s])$$

Passive Temporal Difference Learning

1 TD Learning

- We start with a policy π
- $U(s)$ — Utility for s (init. null)
- $N(s)$ — # times we visited s
- α — Learning rate,
e.g. $\frac{1}{N[s]+1}$

	1	2	3	4
a	0	0	0	+1
b	0			-1
c				

Algorithm consists of:


- Follow π from the initial state to an end state (multiple times);
- Keep track of the number of times we visited each state;
- Update the utility of the states as we go, using the following algorithm:
if s' is new **then** $U[s'] \leftarrow r'$
if s is not null **then**
 increment $N_s[s]$

$$U[s] \leftarrow U[s] + \alpha(N_s[s])(r + \gamma U[s'] - U[s])$$

Passive Temporal Difference Learning

1 TD Learning

- We start with a policy π
- $U(s)$ — Utility for s (init. null)
- $N(s)$ — # times we visited s
- α — Learning rate,
e.g. $\frac{1}{N[s]+1}$



Algorithm consists of:

- Follow π from the initial state to an end state (multiple times);
- Keep track of the number of times we visited each state;
- Update the utility of the states as we go, using the following algorithm:

if s' is new **then** $U[s'] \leftarrow r'$

if s is not null **then**

increment $N_s[s]$


$U[s] \leftarrow U[s] + \alpha(N_s[s])(r + \gamma U[s'] - U[s])$

$$U[a3] \leftarrow 0 + 1/2(0 + 1 - 0) - \text{for } \gamma = 1$$

Passive Temporal Difference Learning

1 TD Learning

- We start with a policy π
- $U(s)$ — Utility for s (init. null)
- $N(s)$ — # times we visited s
- α — Learning rate,
e.g. $\frac{1}{N[s]+1}$



Algorithm consists of:

- Follow π from the initial state to an end state (multiple times);
- Keep track of the number of times we visited each state;
- Update the utility of the states as we go, using the following algorithm:

if s' is new **then** $U[s'] \leftarrow r'$

if s is not null **then**

increment $N_s[s]$


$$U[s] \leftarrow U[s] + \alpha(N_s[s])(r + \gamma U[s'] - U[s])$$

$$U[a2] \leftarrow 0 + 1/3(0 + 1/2 - 0) - \text{for } \gamma = 1$$

Passive Temporal Difference Learning

1 TD Learning

- We start with a policy π
- $U(s)$ — Utility for s (init. null)
- $N(s)$ — # times we visited s
- α — Learning rate,
e.g. $\frac{1}{N[s]+1}$



Algorithm consists of:

- Follow π from the initial state to an end state (multiple times);
- Keep track of the number of times we visited each state;
- Update the utility of the states as we go, using the following algorithm:

if s' is new **then** $U[s'] \leftarrow r'$

if s is not null **then**

increment $N_s[s]$

$$U[s] \leftarrow U[s] + \alpha(N_s[s])(r + \gamma U[s'] - U[s])$$

$$U[a3] \leftarrow 1/2 + 1/3(0 + 1 - 1/2) - \text{for } \gamma = 1$$

Passive TD Algorithm

1 TD Learning

```


function PASSIVE-TD-AGENT(percept) returns an action
  inputs: percept, a percept indicating the current state  $s'$  and reward signal  $r'$ 
  persistent:  $\pi$ , a fixed policy
     $U$ , a table of utilities, initially empty
     $N_s$ , a table of frequencies for states, initially zero
     $s, a, r$ , the previous state, action, and reward, initially null

  if  $s'$  is new then  $U[s'] \leftarrow r'$ 
  if  $s$  is not null then
    increment  $N_s[s]$ 
     $U[s] \leftarrow U[s] + \alpha(N_s[s])(r + \gamma U[s'] - U[s])$ 
  if  $s'.\text{TERMINAL?}$  then  $s, a, r \leftarrow \text{null}$  else  $s, a, r \leftarrow s', \pi[s'], r'$ 
  return  $a$ 


```

Passive Agent Results

1 TD Learning




(a)



(b)

Problems with Passive RL

1 TD Learning



T F

Long Convergence


Limited by Policy

Missing States

Poor Estimate

Problems with Passive RL

1 TD Learning



T	F	
X		Long Convergence
X		Limited by Policy
X		Missing States
X		Poor Estimate

Active Greedy TD

1 TD Learning

- Active reinforcement learning seeks to use the reward information learned by the passive algorithm to generate a new policy

- Active reinforcement learning seeks to use the reward information learned by the passive algorithm to generate a new policy
- **Active Greedy** Temporal Difference learning

Active Greedy TD

1 TD Learning

- Active reinforcement learning seeks to use the reward information learned by the passive algorithm to generate a new policy
- **Active Greedy Temporal Difference learning**
 - Works exactly like passive TD learning, but

Active Greedy TD

1 TD Learning

- Active reinforcement learning seeks to use the reward information learned by the passive algorithm to generate a new policy
- **Active Greedy Temporal Difference learning**
 - Works exactly like passive TD learning, but
 - After a certain number of iterations, it recomputes a new policy π

- Active reinforcement learning seeks to use the reward information learned by the passive algorithm to generate a new policy
- **Active Greedy Temporal Difference learning**
 - Works exactly like passive TD learning, but
 - After a certain number of iterations, it recomputes a new policy $\pi \rightarrow \pi_2$
 - This new policy is computed by **solving the MDP** using the estimated utilities

Active Greedy TD

1 TD Learning

- Active reinforcement learning seeks to use the reward information learned by the passive algorithm to generate a new policy
- **Active Greedy** Temporal Difference learning
 - Works exactly like passive TD learning, but
 - After a certain number of iterations, it recomputes a new policy $\pi \rightarrow \pi_2$
 - This new policy is computed by **solving the MDP** using the estimated utilities

$$\pi(s) = \arg \max_a \sum_{s'} P(s'|s, a) * V(s')$$

- Use the new policy instead of the old one for passive TD

Active Greedy TD

1 TD Learning


- Active reinforcement learning seeks to use the reward information learned by the passive algorithm to generate a new policy
- **Active Greedy** Temporal Difference learning
 - Works exactly like passive TD learning, but
 - After a certain number of iterations, it recomputes a new policy $\pi \rightarrow \pi_2$
 - This new policy is computed by **solving the MDP** using the estimated utilities

$$\pi(s) = \arg \max_a \sum_{s'} P(s'|s, a) * U(s')$$


- Use the new policy instead of the old one for passive TD

Greedy Agent Results

1 TD Learning



(a)




(b)


Is this policy optimal?

Greedy Agent Results

1 TD Learning



(a)




(b)


Is this policy optimal? **No**

Exploration vs. Exploitation

1 TD Learning



(a)




(b)


- Even after changing policy, we are still stuck to something close to the original policy

Exploration vs. Exploitation

1 TD Learning



(a)




(b)


- Even after changing policy, we are still stuck to something close to the original policy
- So we need to balance:

Exploration vs. Exploitation

1 TD Learning



(a)





(b)

- Even after changing policy, we are still stuck to something close to the original policy
- So we need to balance:
 - exploiting the values we discovered

Exploration vs. Exploitation



1 TD Learning



- Even after changing policy, we are still stuck to something close to the original policy
- So we need to balance:
 - exploiting the values we discovered
 - explore different actions to discover new opportunities

Exploration vs. Exploitation


1 TD Learning




- Even after changing policy, we are still stuck to something close to the original policy
- So we need to balance:
 - exploiting the values we discovered
 - explore different actions to discover new opportunities

Exploration vs. Exploitation

1 TD Learning



(a)



(b)

- Even after changing policy, we are still stuck to something close to the original policy
- So we need to balance:
 - exploiting the values we discovered
 - explore different actions to discover new opportunities

one possibility is to choose random actions occasionally

Errors in Utility

1 TD Learning


Possible problems with these methods

- We might not have sampled enough N
- We might have started with a very poor policy π

What can happen if we vary the following

Sampling	Policy		
		T	F

parameters?



Underestimate U
 Overestimate U
 Can we improve
 with higher N


Errors in Utility

1 TD Learning

Possible problems with these methods

- We might not have sampled enough N
- We might have started with a very poor policy π

What can happen if we vary the following



parameters?	Sampling		Policy		Underestimate U
	T	F	T	F	
	X		X		Overestimate U
	X		X		Can we improve with higher N
	X				

Exploration Agents

1 TD Learning

if s' is new **then** $U[s'] \leftarrow r'$

if s is not null **then**

increment $N_s[s]$

$U[s] \leftarrow U[s] + \alpha(N_s[s])(r + \gamma U[s'] - U[s])$

- Mitigate the negative effects of having explored too little using an optimistic estimate function

Exploration Agents

1 TD Learning

if s' is new **then** $U[s'] \leftarrow r'$

if s is not null **then**

increment $N_s[s]$

$U[s] \leftarrow U[s] + \alpha(N_s[s])(r + \gamma U[s'] - U[s])$

- Mitigate the negative effects of having explored too little using an optimistic estimate function

if s' is new **then** $U[s'] \leftarrow r'$

if s is not null **then**

increment $N_s[s]$

$$U[s] \leftarrow U[s] + \alpha(N_s[s])(r + \gamma U[s'] - U[s])$$

- Mitigate the negative effects of having explored too little using an optimistic estimate function

$$f(u, n) = \begin{cases} R^+ & \text{if } n < N_e \\ u & \text{otherwise} \end{cases}$$

where:

- N_e is an exploration threshold, and
- R^+ is the best reward we expect to receive

- With this update, we avoid underestimating a state until we have explored the domain sufficiently

if s' is new **then** $U[s'] \leftarrow r'$

if s is not null **then**

increment $N_s[s]$

$$U[s] \leftarrow U[s] + \alpha(N_s[s])(r + \gamma U[s'] - f(U[s], N_s[s]))$$

- Mitigate the negative effects of having explored too little using an optimistic estimate function

$$f(u, n) = \begin{cases} R^+ & \text{if } n < N_e \\ u & \text{otherwise} \end{cases}$$


where:

- N_e is an exploration threshold, and
- R^+ is the best reward we expect to receive


- With this update, we avoid underestimating a state until we have explored the domain sufficiently

Exploration Agent Result

1 TD Learning



(a)




(b)



Outline


2 Q Learning

- ▶ TD Learning
- ▶ Q Learning
- ▶ Feature Generalization



- Once we have learned the utility of the states, we simply apply the Bellman equation to select the best policy:


$$\pi^*(s) = \arg \max_a \sum_{s'} P(s'|s, a) * U(s')$$



- Once we have learned the utility of the states, we simply apply the Bellman equation to select the best policy:

$$\pi^*(s) = \arg \max_a \sum_{s'} P(s'|s, a) * U(s')$$


- But what if we do not know the transition model $P(s'|s, a)$?




- Once we have learned the utility of the states, we simply apply the Bellman equation to select the best policy:

$$\pi^*(s) = \arg \max_a Q(s, a)$$

- But what if we do not know the transition model $P(s'|s, a)$?
- We can use a method called Q learning, that learns a different value $Q(s, a)$, from which we can derive the optimal policy




- Instead of storing rewards for each state, we store rewards for each action we took at each state
this is Q value $Q(s, a)$




- Instead of storing rewards for each state, we store rewards for each action we took at each state
this is Q value $Q(s, a)$
- The update algorithm is quite similar to TD, but with the following update function:

$$Q(s, a) \leftarrow Q(s, a) + \alpha(R(s) + \gamma \max_{a'} Q(s', a') - Q(s, a))$$




- Instead of storing rewards for each state, we store rewards for each action we took at each state
this is Q value $Q(s, a)$
- The update algorithm is quite similar to TD, but with the following update function:

$$Q(s, a) \leftarrow Q(s, a) + \alpha(R(s) + \gamma \max_{a'} Q(s', a') - Q(s, a))$$



- Instead of storing rewards for each state, we store rewards for each action we took at each state
this is Q value $Q(s, a)$
- The update algorithm is quite similar to TD, but with the following update function:

$$Q(s, a) \leftarrow Q(s, a) + \alpha(R(s) + \gamma \max_{a'} Q(s', a') - Q(s, a))$$



- Instead of storing rewards for each state, we store rewards for each action we took at each state
this is Q value $Q(s, a)$
- The update algorithm is quite similar to TD, but with the following update function:

$$Q(s, a) \leftarrow Q(s, a) + \alpha(R(s) + \gamma \max_{a'} Q(s', a') - Q(s, a))$$

Q Learning Algorithm

2 Q Learning

function Q-LEARNING-AGENT(*percept*) **returns** an action

inputs: *percept*, a percept indicating the current state s' and reward signal r'

persistent: Q , a table of action values indexed by state and action, initially zero

N_{sa} , a table of frequencies for state-action pairs, initially zero

s, a, r , the previous state, action, and reward, initially null

if TERMINAL?(*s*) **then** $Q[s, \text{None}] \leftarrow r'$

if *s* is not null **then**

increment $N_{sa}[s, a]$

$Q[s, a] \leftarrow Q[s, a] + \alpha(N_{sa}[s, a])(r + \gamma \max_{a'} Q[s', a'] - Q[s, a])$

$s, a, r \leftarrow s', \text{argmax}_{a'} f(Q[s', a'], N_{sa}[s', a']), r'$

return *a*

- Exploratory Q-learning agent: same exploration function as TD-learning
- Computes the best action at each call through arg max

- SARSA – State-Action-Reward-State-Action is a close relative to Q-Learning
- Algorithm is exactly the same, but the update rule is slightly different

$$Q(s, a) \leftarrow Q(s, a) + \alpha(R(s) + \gamma Q(s', a') - Q(s, a))$$

- Where a' – action actually taken in s' (notice max is gone)
- Subtle differences:

- SARSA – State-Action-Reward-State-Action is a close relative to Q-Learning
- Algorithm is exactly the same, but the update rule is slightly different

$$Q(s, a) \leftarrow Q(s, a) + \alpha(R(s) + \gamma Q(s', a') - Q(s, a))$$

- Where a' – action actually taken in s' (notice max is gone)
- Subtle differences:
 - Q-learning backs up the best action (**off-policy**)

- SARSA – State-Action-Reward-State-Action is a close relative to Q-Learning
- Algorithm is exactly the same, but the update rule is slightly different

$$Q(s, a) \leftarrow Q(s, a) + \alpha(R(s) + \gamma Q(s', a') - Q(s, a))$$

- Where a' – action actually taken in s' (notice max is gone)
- Subtle differences:
 - Q-learning backs up the best action (**off-policy**)
 - SARSA backs up the action actually taken (**on-policy**)

- SARSA – State-Action-Reward-State-Action is a close relative to Q-Learning
- Algorithm is exactly the same, but the update rule is slightly different

$$Q(s, a) \leftarrow Q(s, a) + \alpha(R(s) + \gamma Q(s', a') - Q(s, a))$$

- Where a' – action actually taken in s' (notice max is gone)
- Subtle differences:
 - Q-learning backs up the best action (**off-policy**)
 - SARSA backs up the action actually taken (**on-policy**)
 - Algorithms are **identical** when there is **no exploration**



Outline

3 Feature Generalization

- ▶ TD Learning
- ▶ Q Learning
- ▶ Feature Generalization

Generalization

3 Feature Generalization

- TD-RL and Q-learning use a monolithic lookup table
 - Will work for 2D maze-like environments, and up to ≈ 10000 states
 - Not good enough for, e.g. chess, backgammon ($10^{20} - 10^{40}$ states)

- TD-RL and Q-learning use a monolithic lookup table
 - Will work for 2D maze-like environments, and up to ≈ 10000 states
 - Not good enough for, e.g. chess, backgammon ($10^{20} - 10^{40}$ states)
- We could generalize using **function approximation**
 - Use an alternative representation for the Q-function, e.g. a linear function of features:

$$\hat{U}_\theta(s) = \theta_1 f_1(s) + \theta_2 f_2(s) + \dots + \theta_n f_n(s)$$

- For direct utility estimation: exactly like **supervised learning**

- TD-RL and Q-learning use a monolithic lookup table
 - Will work for 2D maze-like environments, and up to ≈ 10000 states
 - Not good enough for, e.g. chess, backgammon ($10^{20} - 10^{40}$ states)
- We could generalize using **function approximation**
 - Use an alternative representation for the Q-function, e.g. a linear function of features:
$$\hat{U}_\theta(s) = \theta_1 f_1(s) + \theta_2 f_2(s) + \dots + \theta_n f_n(s)$$
 - For direct utility estimation: exactly like **supervised learning**
- More recently, use a deep neural network to approximate Q
 - Not exactly like supervised learning

Generalization with Features

3 Feature Generalization

- For example, consider the utilities for our 4×3 grid, its only features are the x and y coordinates, so:

$$\hat{U}_\theta(x, y) = \theta_0 + \theta_1 x + \theta_2 y$$

Generalization with Features

3 Feature Generalization

- For example, consider the utilities for our 4×3 grid, its only features are the x and y coordinates, so:

$$\hat{U}_\theta(x, y) = \theta_0 + \theta_1 x + \theta_2 y$$

If $(\theta_0, \theta_1, \theta_2) = (0.5, 0.2, 0.1)$, then $\hat{U}_\theta(1, 1) = 0.8$

Generalization with Features

3 Feature Generalization

- For example, consider the utilities for our 4×3 grid, its only features are the x and y coordinates, so:

$$\hat{U}_\theta(x, y) = \theta_0 + \theta_1 x + \theta_2 y$$

If $(\theta_0, \theta_1, \theta_2) = (0.5, 0.2, 0.1)$, then $\hat{U}_\theta(1, 1) = 0.8$

Suppose we run a trial and obtain $\hat{U}_\theta(1, 1) = 0.4$,
then we need to correct for error

Generalization with Features

3 Feature Generalization

- For example, consider the utilities for our 4×3 grid, its only features are the x and y coordinates, so:

$$\hat{U}_\theta(x, y) = \theta_0 + \theta_1 x + \theta_2 y$$

If $(\theta_0, \theta_1, \theta_2) = (0.5, 0.2, 0.1)$, then $\hat{U}_\theta(1, 1) = 0.8$

Suppose we run a trial and obtain $\hat{U}_\theta(1, 1) = 0.4$,
then we need to correct for error

- Let $u_j(s)$ be the observed utility of s at trial j
With an error function $E_j(s) = (\hat{U}_\theta(s) - u_j(s))^2/2$, we can move each parameter θ_i following the rate of change defined by $\delta E_j / \delta \theta_i$, using:

$$\theta_i \leftarrow \theta_i - \alpha \frac{\delta E_j}{\delta \theta_i} = \theta_i - \alpha(u_j(s) - \hat{U}_\theta(s)) \frac{\delta \hat{U}_\theta(s)}{\delta \theta_i}$$

Generalization with Features (Update Rules)

3 Feature Generalization

$$\theta_i \leftarrow \theta_i - \alpha(u_j(s) - \hat{U}_\theta(s)) \frac{\delta \hat{U}_\theta(s)}{\delta \theta_i}$$

- This is called the Widrow-Hoff rule or delta rule for online least squares
- If we apply this rule to our linear function $\hat{U}_\theta(x, y) = \theta_0 + \theta_1 x + \theta_2 y$, we obtain three simple update rules

$$\theta_0 \leftarrow \theta_0 + \alpha(u_j(s) - \hat{U}_\theta(s)),$$

$$\theta_1 \leftarrow \theta_1 + \alpha(u_j(s) - \hat{U}_\theta(s))x,$$

$$\theta_2 \leftarrow \theta_2 + \alpha(u_j(s) - \hat{U}_\theta(s))y.$$

- So, for every transition we make in a single state, we update the utilities of **every other state**

Generalization with Features (TD Updates)

3 Feature Generalization

- Given a collection of samples, we can estimate the utility of some states without ever visiting them
- We can apply this same principle for TD-RL
- And for Q-learning:

Generalization with Features (TD Updates)

3 Feature Generalization

- Given a collection of samples, we can estimate the utility of some states without ever visiting them
- We can apply this same principle for TD-RL

$$\theta_i \leftarrow \theta_i + \alpha \left[R(s) + \gamma \hat{U}_\theta(s') - \hat{U}_\theta(s) \right] \frac{\partial \hat{U}_\theta(s)}{\partial \theta_i}$$

- And for Q-learning:

Generalization with Features (TD Updates)

3 Feature Generalization

- Given a collection of samples, we can estimate the utility of some states without ever visiting them
- We can apply this same principle for TD-RL

$$\theta_i \leftarrow \theta_i + \alpha \left[R(s) + \gamma \hat{U}_\theta(s') - \hat{U}_\theta(s) \right] \frac{\partial \hat{U}_\theta(s)}{\partial \theta_i}$$

- And for Q-learning:

$$\theta_i \leftarrow \theta_i + \alpha \left[R(s) + \gamma \max \hat{Q}_\theta(s', a') - \hat{Q}_\theta(s, a) \right] \frac{\partial \hat{Q}_\theta(s, a)}{\partial \theta_i}$$

Convergence of Prediction Algorithms

3 Feature Generalization

On/Off-Policy	Algorithm	Table Lookup	Linear	Non-Linear
On-Policy	MC	✓	✓	✓
	TD(α)	✓	✓	✗
	TD(λ)	✓	✓	✗
Off-Policy	MC	✓	✓	✓
	TD(α)	✓	✗	✗
	TD(λ)	✓	✗	✗

Convergence of Prediction Algorithms

3 Feature Generalization

- TD does not follow the gradient of any objective function
- This is why TD can diverge when off-policy or using non-linear function approximation
- **Gradient TD** follows true gradient of projected Bellman error

On/Off-Policy	Algorithm	Table Lookup	Linear	Non-Linear
On-Policy	MC	✓	✓	✓
	TD	✓	✓	✗
	Gradient TD	✓	✓	✓
Off-Policy	MC	✓	✓	✓
	TD	✓	✗	✗
	Gradient TD	✓	✓	✓

Convergence of Control Algorithms (Active)

3 Feature Generalization

Algorithm	Table Lookup	Linear	Non-Linear
MC-Control	✓	(✓)	✗
SARSA	✓	(✓)	✗
Q-Learning	✓	✗	✗
Gradient Q-Learning	✓	✓	✗


(✓) = chatters around near-optimal value function

Deep Q-Networks

3 Feature Generalization

Deep Q-Networks (DQN) was one of the first such NN-based function approximations proven at scale:

Mnih, Volodymyr, et al. **Human-level control through deep reinforcement learning.**
Nature 518.7540 (2015): 529.



Reinforcement Learning Summary

3 Feature Generalization

- How to solve an MDP without P and R, using interaction
 - TD-learning
 - Q-Learning (a family of algorithms)
- The balance between exploration and exploitation
 - Idea: instead of exploration function, use $\hat{Q} \leftarrow \text{awesome}$
 - Start really optimistic and exploration will only drive values down
- How to learn the utility of individual features



UNIVERSITY OF
ABERDEEN



Any Questions.